Evolution of cooperation within a behavior-based perspective: confronting nature and animats*

Samuel Delepoulle[1,2] delepoulle@univ-lille3.fr,
Philippe Preux[2] preux@lil.univ-littoral.fr,
Jean-Claude Darcheville[1] darcheville@univ-lille3.fr

[1] Unité de Recherche sur l'Evolution du Comportement et des Apprentissages, Université de Lille 3, B.P. 149, 59653 Villeneuve d'Ascq Cedex, France.
[2] Laboratoire d'Informatique du Littoral, B.P. 719, 62228 Calais, France.

Abstract. We study the evolution of social behaviors within a behavioral framework. To this end, we define a "minimal social situation" that is experimented with both humans and simulations based on reinforcement learning algorithms. We analyse the dynamics of behaviors in this situation by way of operant conditioning. We show that the best reinforcement algorithm, based on Staddon-Zhang's equations, has a performance and a variety of behaviors that comes close to that of humans, and clearly outperforms the well-known Q-learning. Though we use here a rather simple, yet rich, situation, we argue that operant conditioning deserves much study in the realm of artificial life, being too often misunderstood, and confused with classical conditioning.

1 Motivations: Artificial Life and the experimental analysis of behavior

As stated at its origins, artificial life deals with the study of life as it exists, and life as it might exist [17]. This paper deals with the first part of this project, the study of life as it exists. Our endeavor concerns the study of the dynamics of behavior relying on simple, though sound, nature grounded, assumptions. These assumptions are drawn from a selectionist approach of behaviors, compatible with, and complementary to, the selectionist approach of the evolution of living species. Basically, relying on Skinner's work, the selectionist approach of behaviors (or radical behaviorism, operant or instrumental conditioning [23, 22, 25]), states that a behavior is more likely to be re-executed, and spread in the population, if it is followed by positive consequences. Conversely, a behavior is likely to disappear if it is followed by negative consequences. Furthermore, behaviors being generally never redone twice exactly identically, variation naturally occurs in behaviors, hence a mutation-like mechanism is naturally included. This has much to share with the idea of genes being retained and spreading in a population if it is well adapted [24]. Operant learning is well known to be adequate

* This research is supported by "Conseil Regional Nord-Pas de Calais" (contract n 97 53 0283)

to explain complex behaviors of animals, and to allow animals acquire complex behavior [25]. These behaviors are indeed much more complex than those currently exhibited by top of the art robots [5]. In [7], we have shown that a task consisting in sharing a task between agents can be solved quasi-optimally by a set of agents which behavior evolution is selectionist. In this paper, we focus on the emergence of cooperation among a set of agents and explain it by mean of a selection of behaviors due to the context in which they are emitted, or "contingencies of reinforcement". Our approach is original is many respects. First, we try to keep a pure behavior-based approach, relying on the work of the experimental analysis of behavior. Second, we work in the same time on living beings, mostly human beings, and computer simulations.

2 Evolution of cooperation

What makes the evolution of cooperation possible between two, or more, living beings? In natural situations, one living being frequently behaves in such a way that produces favorable consequences to other living beings. For the individual who emits such a behavior, the immediate payoff can be inexistent, or negative. Hence, the individual does not directly benefit from his, or her, own behaviors. How can such behaviors appear? How can they be retained? These are serious problems arising if we adopt a selectionist way of thinking, such as within the theory of natural selection, or the selection of behaviors by way of their consequences.

In his famous work [2], Axelrod uses the problem of the Iterated Prisoner's Dilemma (IPD) to formulate the problem of cooperation in the framework of game theory. He demonstrates the importance of the repetition of cooperation situations. Dawkins [6] studies the role of genes in cooperative situations. Many other attempts to explain cooperative behaviors have been made [15, 11]. However, none of them have yet tackled the question of the evolution of cooperation within a strict behavioral perspective. Indeed, cooperation is a particular behavior. In this paper, we study the conditions within which cooperation can appear and get installed.

We have worked along a line that confrontates the dynamics of behavior of human beings with that of animats merely implementing reinforcement algorithms. We put a fundamental emphasis on this confrontation for different reasons. To be clear, we obviously do not have the ambition to simulate a human being (!): the algorithms that we use are much cruder. More seriously, for the problem of cooperation in a certain experimental layout, we aim to compare the dynamics of human behaviors with that of reinforcement algorithms. Then, relying on the fact that the algorithms are all selectionist, we suggest that human beings may behave in the same situation according to selectionist principles if both dynamics actually match. Furthermore, this would suggest that a social behavior can be the consequence of individual behaviors. The aim is also to assess the performance of reinforcement algorithms in environments in which we try to carefully keep the very essential features of the real world.

3 A behavioral perspective on human cooperation

3.1 Description of a "minimal social situation"

As stated by Hake and Vuklich [10] within a behavioral framework, cooperation is defined by the fact that the reinforcement of two individuals must be *"at least in part dependent upon the responses of the other individual"*. In cooperative procedures, an individual can improve not only his own payoff but also the payoff of his, or her, party.

This experimental layout is directly inspired by Sidowski's experiments [20, 21]. In this situation, subjects are invited to interact with a computer during thirty minutes. On the computer screen, a window includes two buttons and a counter. When subject A (resp. B) clicks on button 1, B's counter (resp. A's) is incremented (rewarding action, or R action). When A (resp. B) clicks on button 2, B's counter (resp. A's) is decremented (punishing action, or P action). Then, the behavior of a subject has no result on his, or her, own counter but on the counter of his, or her, party. It is noteworthy that a subject may choose to do nothing for a while, or even during the whole experiment. Furthermore, agent actions are not synchronous: they do not have to click at the same time and one click is immediately taken into account to provide its consequence to the other subject. This is a fundamental difference between the situation we use and other published work in the field.

3.2 Results

Thirteen couples of subjects underwent this experiment. During the first minutes of the experiment, clicks are equally shared on the two buttons but after a few minutes the rate of clicks on the button 1 increases. At the end of the experiment, subjects click about five times more on button 1. These results accord with other studies using similar procedures with humans [20, 21], and with animals [4].

3.3 Interpretation

According to the experimental analysis of behavior [23, 25] we can suggest how cooperative behaviors emerge. In this approach, behaviors are selected by their consequences [9]. The behavior of an organism varies; if these variations are profitable to this organism, they are likely to be selected.

In the experiment we made, if the action (R, or P) of both subjects are done synchronously, three combinations are possible.

- If both subjects press button 1, each one gains 1 point. In that case, clicking on button 1 is reinforced and they will continue clicking on button 1,
- If both subjects press button 2, each one loses 1 point. So, they will change their behavior and press button 1. Then, the subjects comes to the previous situation,

– If one subject presses button 1 while the other one presses button 2, the first subject loses 1 point while the second gains 1 point. The first will change his, or her, behavior and the second will not. Then, subjects are in the second situation, which again leads to the first.

B\A	clicks on button 1		clicks on button 2
clicks on button 1	+1 +1		+1 -1
		↖	↓
clicks on button 2	-1 +1	→	-1 -1

Table 1. This table gives the payoff for both subjects of their actions. The first payoff (± 1) of the couple is that of A, the second is that of B. See the text for explanations of the arrows.

Table 1 displays the consequences of the subject's choices if they act synchronously. If their actions are synchronous, the spaces of states is made of 4 states. Arrows show the trajectory among the states. State (+1 +1) is an attractor for all initial states. Thus, cooperation is an attractor for the dynamics of behaviors in this experimental situation: if agents behaves synchronously, cooperation appears very quickly and remains. We observe this phenomenon experimentally: originally behaving asynchronously, cooperation appears in the group exactly at the very moment when responses become synchronized.

Each pair of subjects can be represented in a two dimensional space, the axis being the cooperative rate of the two subjects. Figures 1(a) and 1(c) represent the evolution of the cooperative rate of the 13 couples of subjects. In figure 1(a) representing the begining of the experiment, we notice that reponses form a two dimensional gaussian-like distribution at the center. That distribution may be the result of random responses of the two subjects. Figure 1(c) shows the result during the four last minutes of the experiment. The major part of the distribution is concentrated on a corner, this point represents the maximum cooperation rate for both subjects. Considering the results of each group, we notice that when a pair of subjects is entering in that state, they always stay in it. After a certain amount of time (very variable), most of groups switch to that state.

To finish the analysis of the situation, we notice that a mere stimulus-response architecture is not able to solve the minimal social situation even in this very restricted case where each agent can either reward or punish his, or her, party, and actions are synchronous. A response-stimulus architecture is required which opens the road to operant conditioning.

4 Reinforcement learning algorithm facing cooperation

The interest of the experiment is not to show that human subjects can learn to behave in very simple social situations because we know that humans can

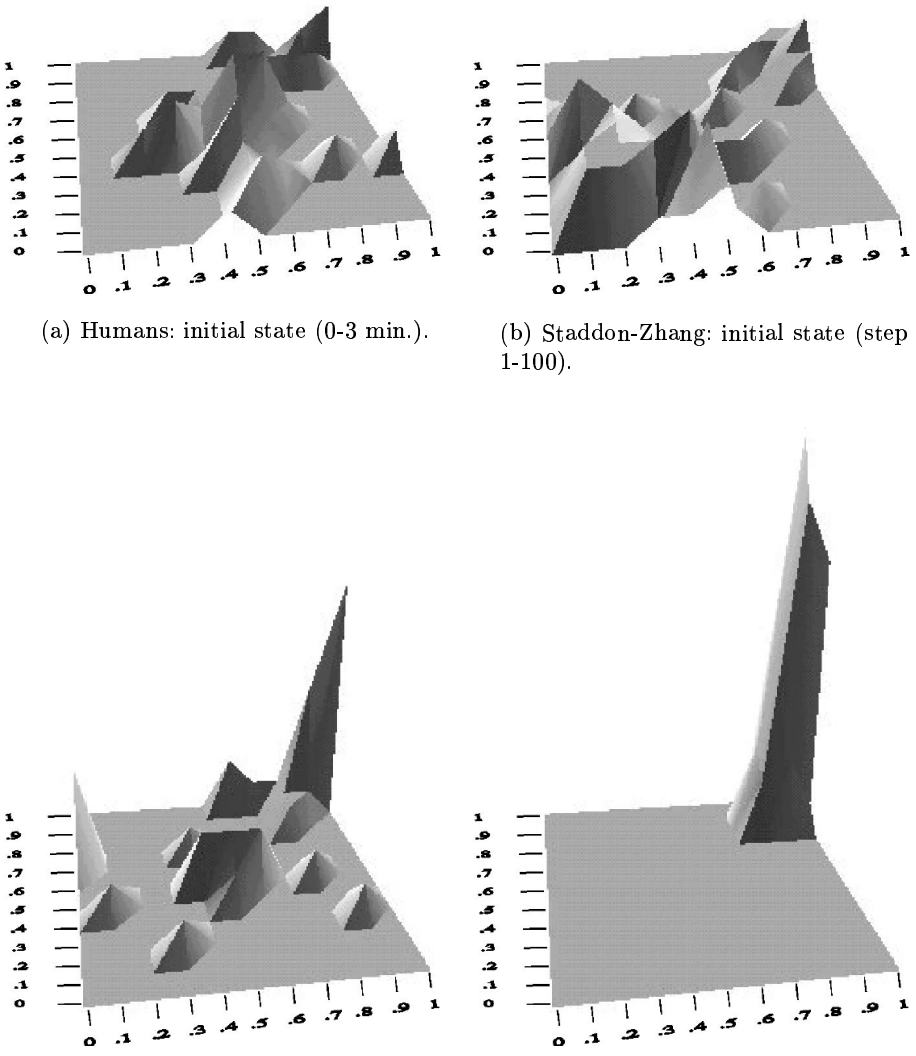

(a) Humans: initial state (0-3 min.).

(b) Staddon-Zhang: initial state (step 1-100).

(c) Humans: final state (26-29 min.).

(d) Staddon-Zhang: final state (step 901-1000).

Fig. 1. Graphical representation of the cooperative rate of the couples of subjects (1^{st} column) and Staddon-Zhang simulation (2^{nd} column) at the beginning of an experiment (a), and at the end of the experiment (c). The abscissa is the rate of clicks on button 1 of subject, or agent, A, the ordinate is the rate of clicks on button 1 of subject, or agent, B. At the beginning of the experiment, the distribution is rather gaussian at the center, meaning that clicks are somehow equally shared between the two buttons. At the end of the experiment, there is a clear spike corresponding to the cooperation: both subjects, or agents, click only on button 1. Of course, some noise always occurs which explains the little blobs here and there.

learn much more complex social situations like imitation or verbal behavior. The interest is to show that a social situation can be explained by the knowledge of the individual law of behavior. Social organization can be the emergent result of individual behaviors. This analysis can also be supported by experimental work on social insects. It was shown that insects living in societies (ants for example) can perform complex tasks without the presence of an individualized central organizer [8, 29]. Many works have used this model to construct self-organised population of agents (see [18, 19] for instance).

The previous section has shown how the behavior of humans subjects can be explained using the principle of behavioral selection. So we propose to use agents based on a selectionist architecture. Reinforcement learning algorithms (RL) meet this requirement [16].

An agent can emit only three behaviors. It can reinforce its party (R), punish its party (P), or do nothing (N). We introduce the "do nothing" behavior to have an asynchronized situation at the beginning of the simulation. As we have shown in the experimental section, synchronization makes selection of cooperation easier because the consequence of a behavior is received immediately.

In the sequel of this section, we present the 5 reinforcement learning algorithms very briefly and give the experimental results in the next section.

4.1 The law of effect

At the end of the XIX$^{\text{th}}$ century, Thorndike [30, 31] has studied the animal behavior and he has suggested a law to predict the evolution of behaviors with regards to their consequences (positive or negative). We formalize his "law of effect" as follows:

$$\text{if } C_i \neq 0, \text{let } s = \frac{C_i.\alpha}{|C_i|}, \text{then} \begin{cases} p_i = \frac{p_i+s}{1+s} \\ p_j = \frac{p_j}{1+s} \quad \text{for } j \neq i \end{cases}$$

where p_i is the probability of apparition of a behavior i, C_i, the consequence of behavior i and α is the learning rate, that is the relative weight of current stimuli with regards to behaviors emitted previously.

4.2 Hilgard and Bower's law

In Hilgard and Bower's law [12], very similar to the law of effect, named "linear reward–inaction algorithm", all non-reinforced actions are weakened. This algorithm always converges with a propability 1 on a particular action (but not always on the best action). This can be expressed by:

$$\text{if } C_i > 0 \begin{cases} p_i = p_i + \alpha(1 - p_i) \\ p_j = p_j - \alpha p_j \qquad \text{for } j \neq i \end{cases}$$

with p_i, C_i and α defined as in the law of effect.

4.3 Staddon-Zhang Model

In 1991, Staddon and Zhang [26] proposed a model in order to solve the *assignment-of-credit* problem without teacher (unsupervised learning). Staddon and Zhang show that this model accounts for qualitative properties of response selection. Their model accords itself not only with regular data in behavior analysis but also with "abnormal" behaviors (autoshaping, superstition and instinctive drift).

In this model, each behavior has a value V_i. All values V_i are in competition and a "winner-take-all" rule is used to select the behavior to emit at each time slot. The values V_i are updated at each time slot according to the equation:

$$V_i = \alpha V_i + \epsilon(1-\alpha) + \beta V_i$$

where $0 < \alpha < 1$ (α is a kind of short term memory parameter), β, the reinforcement parameter should be positive for rewards and negative for punishments, and ϵ is a white noise.

4.4 Action-value method

The goal of this method is to estimate the mean consequence for each behavior and to choose the best one to emit in order to optimize the reward. Sutton [28] gives an iterative method to calculate this estimation. V_i is the estimation of mean consequence and N_i is the number of occurences of behavior i in the past. If behavior i is emitted, then

$$\begin{cases} V_{i+1} = \frac{1}{N_i}[C_i + (N_i - 1).V_i] \\ \\ N_{i+1} = N_i + 1 \end{cases}$$

4.5 Q-Learning

Q-Learning is one of the most famous reinforcement learning method. Based on Sutton and Barto's Time Derivative (TD) model [27], this method has been proposed by Watkins [32]. Q-Learning is an algorithm for solving rapidly and easily stochastic optimal control problem. $Q_{s,a}$ represents the expected future payoff for action a in state s. Q-Learning works by modifing $Q_{s,a}$ for each pair of state-action using the following equation:

$$Q_{s,a} = Q_{s,a} + \alpha[r + \gamma \max_b Q_{s',a} - Q_{s,a}]$$

5 Results of simulation

Agents are tested in two situations. For the first(individual situation), they can choose between three behaviors which have direct consequences on their own counter. In the second situation (interactive situation), agents are put by couples

	individual			social		
algorithm	R	P	N	R	P	N
law of effect	100	0	0	53	16	31
Hilgard and Bower	100	0	0	35	28	37
Mean earning	89	7	4	54	42	4
Q-Learning	89	7	4	58	17	25
Staddon-Zhang	66	0	34	86	1	13

Table 2. Distribution of different behaviors among a hundred agents after one thousand steps of simulation. R represents the behavior that provides reinforcements to its party, P is the bahavior that provides punishment and N is doing nothing. Reinforcements can be given to itself (individual situation) or to the other agent (social situation).

in the same condition than human subjects. The algorithm performs a number of behavior that is rather similar as for human subject

In the individual situation (see columns 2, 3, and 4 of table 2), the environment is very simple and all algorithms succeed in adapting their behavior to optimize their rewards. The difference between algorithms is the speed of convergence. The only algorithm which exhibits a different behavior is Staddon-Zhang's law: within about thirty percent cases, it has a behavior which yields no consequence at all, which may seem a major weakness. However, a careful examination of the behavior of this algorithm shows that it keeps exploring its environment while all other algorithms have settled into a rest point.

In the "social" situation (see columns 5, 6, and 7 of table 2), results are opposite. By far, results of Staddon-Zhang are the best: at the end, it cooperates nearly nine times out of ten. Only Staddon-Zhang, the law of effect, and Q-Learning exhibit results that differs significantly from a random behavior. The very good results of the Staddon-Zhang's algorithm clearly show that the performance of a method designed for stationary problems can be very different in a dynamical situation. As said before, amongst the 5 algorithms, the algorithm relying on Staddon-Zhang's law has the unique feature to keep exploring its environment. This feature might be held as a weakness in a stationary environment, but it is the source of its strength in a dynamical situation. Cooperative rate are plotted on figures 1(b) and 1(d) at the beginning and at the end of simulation from random behavior, the algorithm comes to mostly cooperate.

6 Discussion and perspectives

By using a minimal social situation, we show how cooperation, is possible if the behavior of subjects is selected by their environment. The interest of such situations is twofolds. First, there can be a precise recording of emitted behaviors so that the analysis is possible. Second, they can be simulated by simple adaptive agents.

The minimal social situation used has shown that the asynchronization of agents' behavior is an important issue since synchronization implies cooperation

straightforwardly with a high probability. So, the emergence of the synchronization of the agents' behavior is an important part of the route towards cooperation. More generally, it should be clear that the synchronization of actions, which generally lies implicitly in the background of lots of works dealing with simulations, is not a mere detail.

By using a *minimal social situation*, we may suppose that behavioral selection can be important to account for the emergence of cooperation. The use of adaptive agents in the same situation supports such an analysis. So, adaptive agents are a precious tool to test hypothesis about social behaviors. However, we emphasize that simulations never proove anything firmly. Reciprocally, experiments based on human or animal behavior provide insights into the design of agents. In many cases, animal behavior, as the result of million years of evolution, is able to optimize very complex situations. By knowing the mathematical relation between physical characteristics of environment and behaviors, we might be able to build a new generation of architectures of artificial agents, able to learn in many situation. This kind of agents should be able, for instance, to learn to imitate and to learn from verbal instruction [13, 14].

This procedure — controlled social situation and simulation by agents — will be used in order to study more complex behaviors such as work division or sequential decision making in social situations. If the situation is more complex, many reinforcement learning algorithms used in this paper may become unusable because, apart for Q-Learning, they are "context free". Hence their behavior cannot be truly controlled by the characteristics of their environment, unless they are adapted to integrate the context. Due to its performance in a dynamical environment as well as its soundness with regards to the experimental analysis of behavior, we will work towards making Staddon-Zhang's law context sensitive.

References

[1] Attonaty, J.M., Chatelin, M.H., Garcia, F., and Ndiaye S.M., Using extended machine learning and simulation technics to design crop management stategies. *In EFITA First European Conference for Information Technology in Agriculture*, (1997) Copenhagen

[2] Axelrod, R., *The evolution of cooperation*, Basic Book Inc. (1984)

[3] Bergen, D.E., Hahn, J.K., Bock, P., An adaptive approch for reactive actor design, *Proc. European Conference on Artificial Life*, (1997).

[4] Boren, J.J., An experimental social relation between two monkeys, *Journal of the experimental analysis of behavior*, **9** (1966) 691–700.

[5] David S. Touretzky, Lisa M. Saksida, Skinnerbots, *Proc. 4th Int'l Conf. on Simulation of Adaptive Behavior, From Animals to Animats 4*, Maes, Mataric, Meyer, Pollack, Wilson (eds), MIT Press, 1996

[6] Dawkins, R., *The Selfish Gene*, Oxford University Press, Oxford (1976)

[7] Delepoulle S., Preux Ph., and Darcheville J.C., Partage des tâches et apprentissage par renforcement, *Proc. Journées Francophones d'Apprentissage*, (1998), 201–204 (in french)

[8] Deneubourg, J.L., and Goss S., Collective paterns and decision-making, *Ethology Ecology and Evolution*, **1** (1989) 295–311.

[9] Donahoe, J.W., Burgos, J.E., and Palmer, D.C., A selectionist approach to reinforcement, *Journal of the experimental analysis of behavior*, **60** (1993) 17–40.
[10] Hake, D.F., Vukelich, R., Analysis of the control exerted by a complex cooperation procedure, *Journal of the experimental analysis of behavior*, **19** (1973) 3–16.
[11] Hemelrijk, C.K., Cooperation without genes, games or cognition, *Proc. European Conference on Artificial Life*, (1997).
[12] Hilgard, E.R. and Bower, G.H.: *Theories of learning* (fourth edition). Prentice-Hall, Enblewood Cliffs, NJ.
[13] Hutchinson, W.R. Teaching an agent to speak and listen with understanding: Why and how? Proceedings of the Intelligent Information Agents Workshop, CIKM, Baltimore: http://www.cs.umbc.edu/ cikm/iia/submitted/viewing/whutchi.html
[14] Hutchinson, W.R., *The 7G operant behavior toolkit: Software and documentation*, Boulder, CO: Behavior System.
[15] Ito, A, How do selfish agents learn to cooperate? *Proc. Artificial life 5*, Langton, Shimohara (eds), MIT Press, (1996) 185–192.
[16] Kaelbling, L.P., Littman, M.L., and Moore, A.W., Reinforcement learning: a survey, *Journal of Artificial Intelligence Research*, **4** (1996) 237–285
[17] Langton Ch., Artificial life, *Proc. Artificial Life*, Langton (ed), Addison-Wesley, (1987), 1–47
[18] McFarland D., Towards Robot Cooperation, *Proc of the International Conference on Simulation of Adaptive Behavior: From Animals to Animats 3*, (1994) 440–444.
[19] Murciano, A., Millán, J.R., Learning signaling behaviors and specialization in cooperative agents. *Adaptive Behavior*, **5(1)** (1997) 5–28
[20] Sidowski, J.B., Reward and Punishment in a Minimal Social Situation, *Journal of Experimental Psychology*, **55** (1957) 318–326.
[21] Sidowski, J. B., Wyckoff, B., and Tabory, L., The influence of reinforcement and punishment in a minimal socila situation. *Journal of Abnormal Social Psychology*, **52** (1956) 115–119.
[22] Skinner, B.F., *The behavior of organisms*, (1938). Englewood Cliffs, NJ: Prentice Hall.
[23] Skinner, B.F., *Science and human behavior*, (1953) New York: Macmillan.
[24] Skinner B.F., Selection by consequence, *Science*, **213**, 501–514, 1981
[25] Staddon, J.E.R., *Adaptive Behavior and Learning*, (1981) Cambridge University Press.
[26] Staddon, J.E.R., and Zhang, On the Assignment-of-Credit Problem in Operant Learning, *Neural Network model of Conditioning and Action*, M.L. Caumais S. Grossberg (eds), (1991) Laurence Erlbaum : Hillsdale, N V.
[27] Sutton, R.S. and Barto, A.G., "Time-Derivative Models of Pavlovian Reinforcement", *Learning and Computational Neuroscience: Foundations of Adaptive Networks*, M.Gabriel and J.Moore (eds.), (1990) 497–537. MIT Press: ftp://ftp.cs.umass.edu/pub/anw/pub/sutton/sutton-barto90.ps
[28] Sutton, R.S., *Reinforcement Learning*, MIT Press (1998).
[29] Theraulaz G., Pratte M., and Gervet, J., Behavioural profiles in *Polistes dominulus* (Christ) wasp societies: a quantitative study. *Behaviour* **113**, (1990) 223–250.
[30] Thorndike, E.L., Animal Intelligence: An experimental study of the associative process in animals, *Psychology Monographs*, **2** (1898)
[31] Thorndike, E.L., *Animal Intelligence: Experimental studies.* New York : MacMillan.
[32] Watkins C.J.C.H. and Dayan P., Q-Learning. Technical Note. *Machine Learning*, **8(3)**, (1992), 279–292