

Reinforcement learning

Philippe Preux
philippe.preux@univ-lille.fr
SCOOL



This presentation:

<https://philippe-preux.github.io/talks/AISS-Insa-Rouen/AISS.pdf>

Based on my *Reinforcement Learning lecture notes*, in French only:

<https://philippe-preux.github.io/Documents/digest-ar.pdf>.

Reinforcement learning



The roots of reinforcement learning

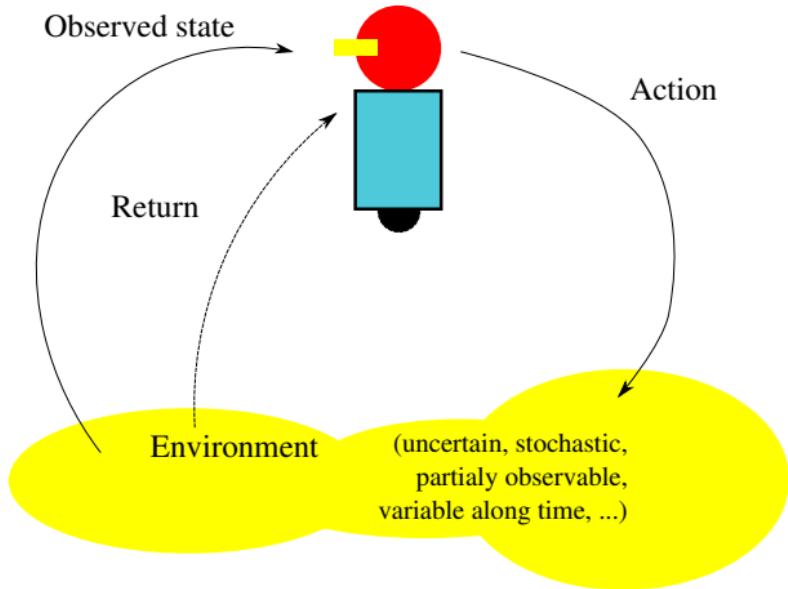
Roots of RL in psychology: the “scientific study of behavior”:

- ▶ Law of effect: Thorndike, 1898.
- ▶ Classical conditioning: Pavlov, 1903, 1927
- ▶ Operant conditioning: Skinner *et al.* from 1931 on.
Key idea: behaviors are selected by their consequences.
(selection of behavior akin selection of species.)
- ▶ Rescorla-Wagner law: 1972.
- ▶ Sutton's Ph.D. defended in Feb. 1984.

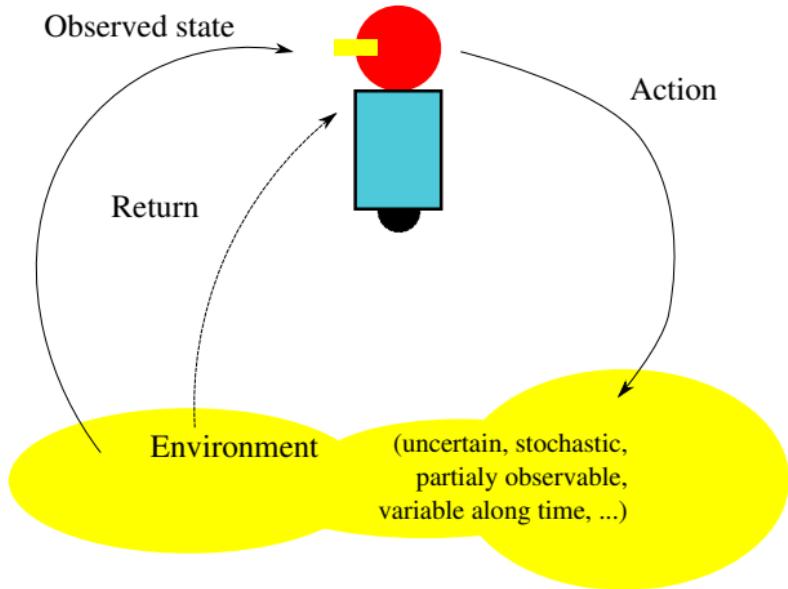
Outline

- ▶ Introduction
- ▶ Markov decision processes and Markov decision problems
- ▶ Reinforcement learning: definition and algorithms
- ▶ RL in practice

Markov decision problems



Markov decision problems



Learn an optimal behavior.

Markov decision process

Markov decision process

describes a dynamical decision system

Definition

A Markov decision process is defined by the tuple $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$ where:

- ▶ \mathcal{T} is the set of instants of decision, $t \in \mathcal{T}$.
For the sake of simplicity, \mathcal{T} is usually the sequence of positive integers: 0, 1, 2, ...

Markov decision process

describes a dynamical decision system

Definition

A Markov decision process is defined by the tuple $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$ where:

- ▶ \mathcal{T} is the set of instants of decision, $t \in \mathcal{T}$.
- ▶ \mathcal{X} is a finite set of states, $x \in \mathcal{X}$.

For the sake of simplicity, the states are usually numbered from 0 to $N - 1$, with $N = |\mathcal{X}|$.

Markov decision process

describes a dynamical decision system

Definition

A Markov decision process is defined by the tuple $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$ where:

- ▶ \mathcal{T} is the set of instants of decision, $t \in \mathcal{T}$.
- ▶ \mathcal{X} is a finite set of states, $x \in \mathcal{X}$.
- ▶ $\mathcal{X}_0 \subset \mathcal{X}$: set of initial states.

Markov decision process

describes a dynamical decision system

Definition

A Markov decision process is defined by the tuple $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$ where:

- ▶ \mathcal{T} is the set of instants of decision, $t \in \mathcal{T}$.
- ▶ \mathcal{X} is a finite set of states, $x \in \mathcal{X}$.
- ▶ $\mathcal{X}_0 \subset \mathcal{X}$: set of initial states.
- ▶ $\mathcal{X}_f \subset \mathcal{X}$: set of terminal states (may be empty).

Markov decision process

describes a dynamical decision system

Definition

A Markov decision process is defined by the tuple $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$ where:

- ▶ \mathcal{T} is the set of instants of decision, $t \in \mathcal{T}$.
- ▶ \mathcal{X} is a finite set of states, $x \in \mathcal{X}$.
- ▶ $\mathcal{X}_0 \subset \mathcal{X}$: set of initial states.
- ▶ $\mathcal{X}_f \subset \mathcal{X}$: set of terminal states (may be empty).
- ▶ \mathcal{A} is a finite set of actions, $a \in \mathcal{A}$.

For the sake of simplicity, actions are usually numbered from 0 to $P - 1$, with $P = |\mathcal{A}|$.

Markov decision process

describes a dynamical decision system

Definition

A Markov decision process is defined by the tuple $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$ where:

- ▶ \mathcal{T} is the set of instants of decision, $t \in \mathcal{T}$.
- ▶ \mathcal{X} is a finite set of states, $x \in \mathcal{X}$.
- ▶ $\mathcal{X}_0 \subset \mathcal{X}$: set of initial states.
- ▶ $\mathcal{X}_f \subset \mathcal{X}$: set of terminal states (may be empty).
- ▶ \mathcal{A} is a finite set of actions, $a \in \mathcal{A}$.
- ▶ \mathcal{P} is the transition function.

Assume at time t , the environment is in state $x_t = x$ and performs action $a_t = a$, then $\mathcal{P}(x, a, x')$ is the probability that the environment will be in state x' at time of decision $t + 1$.

That is: $\mathcal{P}(x, a, x') = Pr[x_{t+1} = x' | x_t = x, a_t = a]$.

Markov decision process

describes a dynamical decision system

Definition

A Markov decision process is defined by the tuple $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$ where:

- ▶ \mathcal{T} is the set of instants of decision, $t \in \mathcal{T}$.
- ▶ \mathcal{X} is a finite set of states, $x \in \mathcal{X}$.
- ▶ $\mathcal{X}_0 \subset \mathcal{X}$: set of initial states.
- ▶ $\mathcal{X}_f \subset \mathcal{X}$: set of terminal states (may be empty).
- ▶ \mathcal{A} is a finite set of actions, $a \in \mathcal{A}$.
- ▶ \mathcal{P} is the transition function.
- ▶ \mathcal{R} is the return function, $r \in \mathbb{R}$.

With the same assumption as for \mathcal{P} : $\mathcal{R}(x, a, x')$ is the expected return for the transition from state x to x' following action a .

That is: $\mathcal{R}(x, a, x') = \mathbb{E}[r_t | x_{t+1} = x', x_t = x, a_t = a]$.

Markov decision process

describes a dynamical decision system

Definition

A Markov decision process is defined by the tuple $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$ where:

- ▶ \mathcal{T} is the set of instants of decision, $t \in \mathcal{T}$.
- ▶ \mathcal{X} is a finite set of states, $x \in \mathcal{X}$.
- ▶ $\mathcal{X}_0 \subset \mathcal{X}$: set of initial states.
- ▶ $\mathcal{X}_f \subset \mathcal{X}$: set of terminal states (may be empty).
- ▶ \mathcal{A} is a finite set of actions, $a \in \mathcal{A}$.
- ▶ \mathcal{P} is the transition function.
- ▶ \mathcal{R} is the return function, $r \in \mathbb{R}$.

Remarks:

Markov decision process

describes a dynamical decision system

Definition

A Markov decision process is defined by the tuple $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$ where:

- ▶ \mathcal{T} is the set of instants of decision, $t \in \mathcal{T}$.
- ▶ \mathcal{X} is a finite set of states, $x \in \mathcal{X}$.
- ▶ $\mathcal{X}_0 \subset \mathcal{X}$: set of initial states.
- ▶ $\mathcal{X}_f \subset \mathcal{X}$: set of terminal states (may be empty).
- ▶ \mathcal{A} is a finite set of actions, $a \in \mathcal{A}$.
- ▶ \mathcal{P} is the transition function.
- ▶ \mathcal{R} is the return function, $r \in \mathbb{R}$.

Remarks:

1. items depend on t and do not depend on $t - 1, t - 2, \dots$: Markov property.

Markov decision process

describes a dynamical decision system

Definition

A Markov decision process is defined by the tuple $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$ where:

- ▶ \mathcal{T} is the set of instants of decision, $t \in \mathcal{T}$.
- ▶ \mathcal{X} is a finite set of states, $x \in \mathcal{X}$.
- ▶ $\mathcal{X}_0 \subset \mathcal{X}$: set of initial states.
- ▶ $\mathcal{X}_f \subset \mathcal{X}$: set of terminal states (may be empty).
- ▶ \mathcal{A} is a finite set of actions, $a \in \mathcal{A}$.
- ▶ \mathcal{P} is the transition function.
- ▶ \mathcal{R} is the return function, $r \in \mathbb{R}$.

Remarks:

1. items depend on t and do not depend on $t - 1, t - 2, \dots$: Markov property.
2. None of these items depend on \mathcal{T} : stationary system (= non autonomous).

Markov decision process

The decision loop

$t \leftarrow 0$

loop

observe state x_t

take action a_t

observe the immediate return r_t

$t \leftarrow t + 1$

end loop

Why would we take an action or an other? What's the point?

Markov decision process

The decision loop

$t \leftarrow 0$

loop

observe state x_t

take action a_t

observe the immediate return r_t

$t \leftarrow t + 1$

end loop

Why would we take an action or an other? What's the point?

An MD process only specifies the dynamics of a system that takes decision: it describes the “how”, not the “why”.

Markov decision process

Example

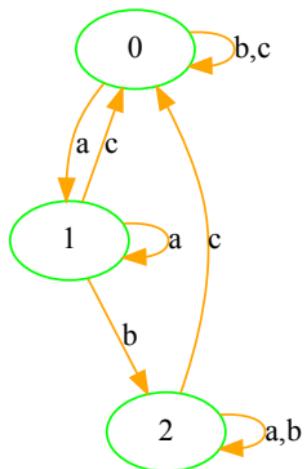
Let's play "21 with a dice".

Rules:

- ▶ you need 1 "standard" dice with 6 faces numbered from 1 to 6.
- ▶ You roll the dice, and note the value on its upper face: that's your initial score.
- ▶ Now repeatedly, you decide whether you roll it again or you stop the game.
- ▶ If you roll the dice, you add the value on its upper face to your current score.

Define a Markov decision process that models this dynamical system.

Running example: my little MDP



a, b, and c are actions.

Transitions are deterministic.

All transitions return 0 except the transition from 2 to 0 that returns 1.

Markov decision problem

The “why”.

Markov decision problem

- ▶ A Markov decision problem describes a sequential decision problem on top of a Markov decision process.

Markov decision problem

- ▶ A Markov decision problem describes a sequential decision problem on top of a Markov decision process.
- ▶ The MD process defines how the dynamical system behaves in reaction to actions.

Markov decision problem

- ▶ A Markov decision problem describes a sequential decision problem on top of a Markov decision process.
- ▶ The MD process defines how the dynamical system behaves in reaction to actions.
- ▶ The MD **problem** defines the **objective of the agent**:
→ we need to define an **objective function**.

Markov decision problem

- ▶ A Markov decision problem describes a sequential decision problem on top of a Markov decision process.
- ▶ The MD process defines how the dynamical system behaves in reaction to actions.
- ▶ The MD **problem** defines the **objective of the agent**:
→ we need to define an objective function.
- ▶ Very often in the RL academic litterature, it is defined as maximizing $\zeta(x_t) = \sum_{k \geq 0} \gamma^k r_{t+k}$, where $\gamma \in [0, 1)$.

Markov decision problem

- ▶ A Markov decision problem describes a sequential decision problem on top of a Markov decision process.
- ▶ The MD process defines how the dynamical system behaves in reaction to actions.
- ▶ The MD **problem** defines the **objective of the agent**:
→ we need to define an objective function.
- ▶ Very often in the RL academic litterature, it is defined as maximizing $\zeta(x_t) = \sum_{k \geq 0} \gamma^k r_{t+k}$, where $\gamma \in [0, 1]$.
- ▶ Solution of an MD problem: a **policy** π that specifies the probability to perform any action a in any state x in order to optimize ζ .
$$\pi(x, a) = Pr[a_t = a | x_t = x], \forall a \in \mathcal{A}.$$

Markov decision problem

- ▶ A Markov decision problem describes a sequential decision problem on top of a Markov decision process.
- ▶ The MD process defines how the dynamical system behaves in reaction to actions.
- ▶ The MD **problem** defines the **objective of the agent**:
→ we need to define an objective function.
- ▶ Very often in the RL academic litterature, it is defined as maximizing $\zeta(x_t) = \sum_{k \geq 0} \gamma^k r_{t+k}$, where $\gamma \in [0, 1)$.
- ▶ Solution of an MD problem: a policy π that specifies the probability to perform any action a in any state x in order to optimize ζ .
 $\pi(x, a) = Pr[a_t = a | x_t = x], \forall a \in \mathcal{A}$.
- ▶ Blackwell's theorem tells us that an optimal policy (for this ζ) is deterministic: $\pi(x)$.
More than 1 action may be optimal in a state.

Markov decision problem

Remarks about the definition

$\gamma?$

Markov decision problem

Remarks about the definition

γ ?

$$\zeta(x_t) = \sum_{k \geq 0} \gamma^k r_{t+k}$$

Markov decision problem

Remarks about the definition

γ ?

$$\zeta(x_t) = \sum_{k \geq 0} \gamma^k r_{t+k}$$

$$\rightarrow \zeta(x_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$

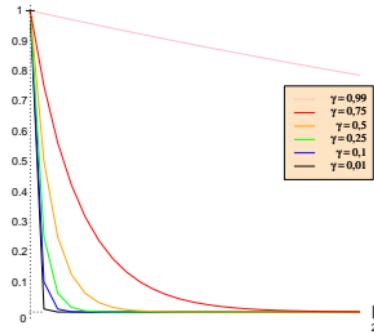
Markov decision problem

Remarks about the definition

γ ?

$$\zeta(x_t) = \sum_{k \geq 0} \gamma^k r_{t+k}$$

$$\rightarrow \zeta(x_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$



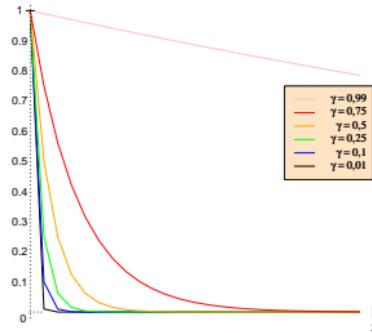
Markov decision problem

Remarks about the definition

$\gamma?$

$$\zeta(x_t) = \sum_{k \geq 0} \gamma^k r_{t+k}$$

$$\rightarrow \zeta(x_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$



- Nice property: If we assume \mathcal{R} is bounded, then ζ converges.

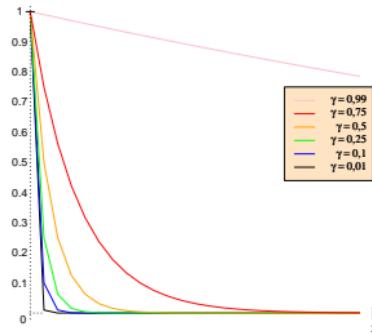
Markov decision problem

Remarks about the definition

$\gamma?$

$$\zeta(x_t) = \sum_{k \geq 0} \gamma^k r_{t+k}$$

$$\rightarrow \zeta(x_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$



- ▶ Nice property: If we assume \mathcal{R} is bounded, then ζ converges.
- ▶ This definition makes the maths easier.

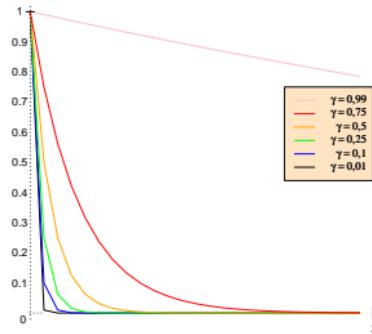
Markov decision problem

Remarks about the definition

$\gamma?$

$$\zeta(x_t) = \sum_{k \geq 0} \gamma^k r_{t+k}$$

$$\rightarrow \zeta(x_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$



- ▶ Nice property: If we assume \mathcal{R} is bounded, then ζ converges.
- ▶ This definition makes the maths easier.
- ▶ Short term consequences vs. longer consequences of actions.

Markov decision problem

A Markov decision process defines the “how”: how does the system evolve in time?

A Markov decision problem defines the “why”: why would the agent choose one or another action?

Answer: to maximize ζ .

Pending questions:

- ▶ Does this problem have a solution?
- ▶ Can we compute it? In practice?
- ▶ How?

Markov decision problem

The value function

- ▶ An “episode” is a sequence of interactions starting from an initial state to a final state.

Markov decision problem

The value function

- ▶ An “episode” is a sequence of interactions starting from an initial state to a final state.
- ▶ A “trajectory” is the set of (x_t, a_t, r_t, x_{t+1}) for all t during an episode.

Markov decision problem

The value function

- ▶ An “episode” is a sequence of interactions starting from an initial state to a final state.
- ▶ A “trajectory” is the set of (x_t, a_t, r_t, x_{t+1}) for all t during an episode.
- ▶ Let $x_0 \in \mathcal{X}_0$.

Markov decision problem

The value function

- ▶ An “episode” is a sequence of interactions starting from an initial state to a final state.
- ▶ A “trajectory” is the set of (x_t, a_t, r_t, x_{t+1}) for all t during an episode.
- ▶ Let $x_0 \in \mathcal{X}_0$.
- ▶ Two episodes starting in x_0 will usually follow different trajectories.

Markov decision problem

The value function

- ▶ An “episode” is a sequence of interactions starting from an initial state to a final state.
- ▶ A “trajectory” is the set of (x_t, a_t, r_t, x_{t+1}) for all t during an episode.
- ▶ Let $x_0 \in \mathcal{X}_0$.
- ▶ Two episodes starting in x_0 will usually follow different trajectories.
- ▶ $\zeta(x), x \in \mathcal{X}_0$ is a random variable.

Markov decision problem

The value function

- ▶ An “episode” is a sequence of interactions starting from an initial state to a final state.
- ▶ A “trajectory” is the set of (x_t, a_t, r_t, x_{t+1}) for all t during an episode.
- ▶ Let $x_0 \in \mathcal{X}_0$.
- ▶ Two episodes starting in x_0 will usually follow different trajectories.
- ▶ $\zeta(x), x \in \mathcal{X}_0$ is a random variable.
- ▶ We define $V^\pi(x)$, the value of a state x according to a policy π by:

$$V^\pi(x) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_{t \geq 0} \sim \pi]$$

the actions being chosen according to policy π .

Markov decision problem

The value function: intuition

$$V^\pi(x) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_{t \geq 0} \sim \pi]$$

$V^\pi(x)$ simply quantifies how good it is to be in state x while behaving according to policy π in order to optimize ζ .

Markov decision problem

The value function: intuition

$$V^\pi(x) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_{t \geq 0} \sim \pi]$$

$V^\pi(x)$ simply quantifies how good it is to be in state x while behaving according to policy π in order to optimize ζ .

We will soon see that given V^π , we can improve the policy and obtain a policy which value is better.

Markov decision problem

Evaluation of the value of a policy

$$V^\pi(x) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_{t \geq 0} \sim \pi]$$

Markov decision problem

Evaluation of the value of a policy

$$V^\pi(x) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_{t \geq 0} \sim \pi]$$



replace ζ by its definition
express the expectation
and you get:

Markov decision problem

Evaluation of the value of a policy

$$V^\pi(x) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_{t \geq 0} \sim \pi]$$



replace ζ by its definition
express the expectation
and you get:

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

Markov decision problem

Evaluation of the value of a policy

$$V^\pi(x) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_{t \geq 0} \sim \pi]$$



replace ζ by its definition
express the expectation
and you get:

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

This may look complicated but all terms are known except the vector V .

Markov decision problem

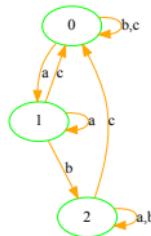
Evaluation of the value of a policy

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

Development → a system of N linear equations with N unknowns, the $V^\pi(x), \forall x \in \mathcal{X}$.

Markov decision problem

Example (continued)



For this MDP , compute the value function of the uniformly random policy, $\pi(x, a) = 1/3$, for $\gamma = 0.9$.

System of linear equations?

Reminder:

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

Markov decision problem

Example (continued)

Compute the value function of the uniformly random policy,
 $\pi(x, a) = 1/3$, for $\gamma = 0.9$.

System of linear equations?

Reminder:

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

Markov decision problem

Example (continued)

Compute the value function of the uniformly random policy,
 $\pi(x, a) = 1/3$, for $\gamma = 0.9$.

System of linear equations?

Reminder:

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

$$\begin{aligned} V^\pi(0) &= \pi(0, a) \sum_{x'} \mathcal{P}(0, a, x') [\mathcal{R}(0, a, x') + \gamma V^\pi(x')] + \\ &\quad \pi(0, b) \sum_{x'} \mathcal{P}(0, b, x') [\mathcal{R}(0, b, x') + \gamma V^\pi(x')] + \\ &\quad \pi(0, c) \sum_{x'} \mathcal{P}(0, c, x') [\mathcal{R}(0, c, x') + \gamma V^\pi(x')] \end{aligned}$$

$$V^\pi(1) = \pi(1, a) \sum_{x'} \mathcal{P}(1, a, x') [\mathcal{R}(1, a, x') + \gamma V^\pi(x')] + \dots$$

$$V^\pi(2) = \pi(2, a) \sum_{x'} \mathcal{P}(2, a, x') [\mathcal{R}(2, a, x') + \gamma V^\pi(x')] + \dots$$

Markov decision problem

Example (continued)

Compute the value function of the uniformly random policy,
 $\pi(x, a) = 1/3$, for $\gamma = 0.9$.

System of linear equations?

Reminder:

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

$$\begin{aligned} V^\pi(0) = & \pi(0, a) \sum_{x'} \mathcal{P}(0, a, x') [\mathcal{R}(0, a, x') + \gamma V^\pi(x')] + \\ & \pi(0, b) \sum_{x'} \mathcal{P}(0, b, x') [\mathcal{R}(0, b, x') + \gamma V^\pi(x')] + \\ & \pi(0, c) \sum_{x'} \mathcal{P}(0, c, x') [\mathcal{R}(0, c, x') + \gamma V^\pi(x')] \dots \end{aligned}$$

→

$$\begin{aligned} V^\pi(0) = & \pi(0, a)(\mathcal{P}(0, a, 0)[\mathcal{R}(0, a, 0) + \gamma V^\pi(0)] + \\ & \mathcal{P}(0, a, 1)[\mathcal{R}(0, a, 1) + \gamma V^\pi(1)] + \\ & \mathcal{P}(0, a, 2)[\mathcal{R}(0, a, 2) + \gamma V^\pi(2)]) + \\ & \pi(0, b)(\mathcal{P}(0, b, 0)[\mathcal{R}(0, b, 0) + \gamma V^\pi(0)] + \\ & \mathcal{P}(0, b, 1)[\mathcal{R}(0, b, 1) + \gamma V^\pi(1)] + \\ & \mathcal{P}(0, b, 2)[\mathcal{R}(0, b, 2) + \gamma V^\pi(2)]) + \\ & \pi(0, c)(\mathcal{P}(0, c, 0)[\mathcal{R}(0, c, 0) + \gamma V^\pi(0)] + \\ & \mathcal{P}(0, c, 1)[\mathcal{R}(0, c, 1) + \gamma V^\pi(1)] + \\ & \mathcal{P}(0, c, 2)[\mathcal{R}(0, c, 2) + \gamma V^\pi(2)]) \end{aligned}$$

$$V^\pi(1) = \dots$$

$$V^\pi(2) = \dots$$

Markov decision problem

Example (continued)

$$\begin{aligned} V^\pi(0) = & \pi(0, a)(\mathcal{P}(0, a, 0)[\mathcal{R}(0, a, 0) + \gamma V^\pi(0)] + \\ & \mathcal{P}(0, a, 1)[\mathcal{R}(0, a, 1) + \gamma V^\pi(1)] + \\ & \mathcal{P}(0, a, 2)[\mathcal{R}(0, a, 2) + \gamma V^\pi(2)]) + \\ & \pi(0, b)(\mathcal{P}(0, b, 0)[\mathcal{R}(0, b, 0) + \gamma V^\pi(0)] + \\ & \mathcal{P}(0, b, 1)[\mathcal{R}(0, b, 1) + \gamma V^\pi(1)] + \\ & \mathcal{P}(0, b, 2)[\mathcal{R}(0, b, 2) + \gamma V^\pi(2)]) + \\ & \pi(0, c)(\mathcal{P}(0, c, 0)[\mathcal{R}(0, c, 0) + \gamma V^\pi(0)] + \\ & \mathcal{P}(0, c, 1)[\mathcal{R}(0, c, 1) + \gamma V^\pi(1)] + \\ & \mathcal{P}(0, c, 2)[\mathcal{R}(0, c, 2) + \gamma V^\pi(2)]) \\ V^\pi(1) = & \dots \\ V^\pi(2) = & \dots \end{aligned}$$

→

$$\begin{aligned} V^\pi(0) = & 1/3(0 \times [...] + \\ & 1 \times [0 + \gamma V^\pi(1)] + \\ & 0 \times [...] + \\ & 1/3(1 \times [0 + \gamma V^\pi(0)] + \\ & 0 \times [...] + \\ & 0 \times [...] + \\ & 1/3(1 \times [0 + \gamma V^\pi(0)] + \\ & 0 \times [...] + \\ & 0 \times [...] + \\ V^\pi(1) = & \dots \\ V^\pi(2) = & \dots \end{aligned}$$

Markov decision problem

Example (continued)

$$\begin{aligned} V^\pi(0) = & \quad 1/3(0 \times [...] + \\ & 1 \times [0 + \gamma V^\pi(1)] + \\ & 0 \times [...] + \\ & 1/3(1 \times [0 + \gamma V^\pi(0)] + \\ & 0 \times [...] + \\ & 0 \times [...] + \\ & 1/3(1 \times [0 + \gamma V^\pi(0)] + \\ & 0 \times [...] + \\ & 0 \times [...] + \end{aligned}$$

$$\begin{aligned} V^\pi(1) = & \quad \dots \\ V^\pi(2) = & \quad \dots \end{aligned}$$

→

$$\begin{aligned} V^\pi(0) = & \quad 1/3(\gamma V^\pi(1) + \gamma V^\pi(0) + \gamma V^\pi(0)) \\ V^\pi(0) = & \quad 1/3(0.9V^\pi(1) + 0.9V^\pi(0) + 0.9V^\pi(0)) \\ \rightarrow & \quad 0.4V^\pi(0) - 0.3V^\pi(1) = 0 \end{aligned}$$

$$\begin{aligned} V^\pi(1) = & \quad \dots \\ V^\pi(2) = & \quad \dots \end{aligned}$$

Markov decision problem

Example (continued)

We obtain the system of linear equations:

$$\begin{pmatrix} 0.4 & -0.3 & 0 \\ -0.3 & 0.7 & -0.3 \\ -0.3 & 0 & 0.4 \end{pmatrix} V^\pi = \begin{pmatrix} 0 \\ 0 \\ 1/3 \end{pmatrix}$$

→

$$V^\pi = \begin{pmatrix} 0,61 \\ 0,82 \\ 1,29 \end{pmatrix}$$

Markov decision problem



Evaluation of the value of a policy:

Implement this approach in a generic way to solve any MDP.

Check your implementation on the little example.

Apply it to the “21 with a dice” problem.

Markov decision problem

Evaluation of the value of a policy: the dynamic programming approach

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

Development → a system of N linear equations with N unknowns,
 $V^\pi(x), \forall x \in \mathcal{X}$.

In principle, an easy problem.

In practice, when N is large (e.g. 10^9), this is a challenging problem.

$V^\pi = \mathbf{A}V^\pi + \mathbf{B}$, where $\mathbf{A} \in \mathbb{R}^{N \times N}$ is an $N \times N$ matrix, and $\mathbf{B} \in \mathbb{R}^N$.
→ $(Id - \mathbf{A})V^\pi = \mathbf{B}$

Thanks to contraction properties of \mathbf{A} , this can be solved iteratively.



express \mathbf{A} and \mathbf{B} in terms of $\pi, \mathcal{P}, \mathcal{R}, \gamma$.

Markov decision problem

Evaluation of the value of a policy: the dynamic programming approach

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

Thanks to contraction properties of \mathbf{A} , this can be solved iteratively.

$$k \leftarrow 0$$

$$V_k \leftarrow 0, \forall x \in \mathcal{X}$$

$$\epsilon \leftarrow 10^{-6} // \text{ some small value}$$

repeat

$$V_{k+1} \leftarrow \dots$$

$$\Delta \leftarrow \|V_k - V_{k+1}\|_\infty$$

$$k \leftarrow k + 1$$

$$\text{until } \Delta < \epsilon \frac{1-\gamma}{2\gamma}$$

This algorithm provides an estimation of V^π at most ϵ away from its true value.

It is very easy to implement and very fast, even when $N = 10^9$.

Markov decision problem

Evaluation of the value of a policy: the dynamic programming approach

$$V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

is known as a “Bellman equation”.

[Bellman, *Dynamic Programming*, Princeton U. Press, 1957]

Markov decision problem

Example (continued): the dynamic programming approach



Implement this dynamic programming approach for policy value evaluation and run it on the two examples.

Check that both methods give the same result.

Markov decision problem

Evaluation of the value of a policy: the Monte Carlo approach



Markov decision problem

Evaluation of the value of a policy: the Monte Carlo approach

Very simple approach: consists in simulating the Markov chain and estimating V^π by computing ζ .

To estimate $V^\pi(x)$:

$cr \leftarrow 0$

$cnt \leftarrow 0$

for $k \in \{0, \dots, K - 1\}$ **do**

$t \leftarrow 0$

$x_t \leftarrow x$

repeat

 sample $a_t \sim \pi(x_t, .)$

 sample the next state $x_{t+1} \sim \mathcal{P}(x_t, a_t, .)$

 sample $r_t \sim \mathcal{R}(x_t, a_t, .)$

$cr \leftarrow cr + \gamma^t r_t$

$cnt \leftarrow cnt + 1$

until γ^t is very small

end for $V^\pi(x) \leftarrow \frac{cr}{cnt}$

Markov decision problem

Evaluation of the value of a policy: the Monte Carlo approach

- ▶ When the MDP is deterministic, the double loop becomes a single loop.
Wrt. traditional tree search algorithms, it is competitive in some cases: very wide tree, or very deep tree, or stochastic dynamics.

Markov decision problem

Evaluation of the value of a policy: the Monte Carlo approach

- ▶ When the MDP is deterministic, the double loop becomes a single loop.
Wrt. traditionnal tree search algorithms, it is competitive in some cases: very wide tree, or very deep tree, or stochastic dynamics.
- ▶ The drawback is that the convergence is slow: $\mathcal{O}(\sqrt{K})$.

Markov decision problem

Evaluation of the value of a policy: the Monte Carlo approach

- ▶ When the MDP is deterministic, the double loop becomes a single loop.

Wrt. traditionnal tree search algorithms, it is competitive in some cases: very wide tree, or very deep tree, or stochastic dynamics.
- ▶ The drawback is that the convergence is slow: $\mathcal{O}(\sqrt{K})$.
- ▶ However, as this is a very simple algorithm, its inner loop is fast.

Markov decision problem

Evaluation of the value of a policy: the Monte Carlo approach

- ▶ When the MDP is deterministic, the double loop becomes a single loop.

Wrt. traditionnal tree search algorithms, it is competitive in some cases: very wide tree, or very deep tree, or stochastic dynamics.

- ▶ The drawback is that the convergence is slow: $\mathcal{O}(\sqrt{K})$.
- ▶ However, as this is a very simple algorithm, its inner loop is fast.
- ▶ One run (inner loop) is named a *rollout*.

Markov decision problem

Evaluation of the value of a policy: the Monte Carlo approach

- ▶ When the MDP is deterministic, the double loop becomes a single loop.

Wrt. traditionnal tree search algorithms, it is competitive in some cases: very wide tree, or very deep tree, or stochastic dynamics.
- ▶ The drawback is that the convergence is slow: $\mathcal{O}(\sqrt{K})$.
- ▶ However, as this is a very simple algorithm, its inner loop is fast.
- ▶ One run (inner loop) is named a *rollout*.
- ▶ Monte Carlo tree search (MCTS) is a more sophisticated version of this basic algorithm to deal with large trees, and stochastic dynamics.

Markov decision problem

Evaluation of the value of a policy: the Monte Carlo approach

- ▶ When the MDP is deterministic, the double loop becomes a single loop.
Wrt. traditionnal tree search algorithms, it is competitive in some cases: very wide tree, or very deep tree, or stochastic dynamics.
- ▶ The drawback is that the convergence is slow: $\mathcal{O}(\sqrt{K})$.
- ▶ However, as this is a very simple algorithm, its inner loop is fast.
- ▶ One run (inner loop) is named a *rollout*.
- ▶ Monte Carlo tree search (MCTS) is a more sophisticated version of this basic algorithm to deal with large trees, and stochastic dynamics.



Implement a Monte Carlo algorithm for policy evaluation. Apply it to the same examples as above and check that the 3 approaches provide the same results.

Markov decision problem

Policy improvement

- ▶ Once the value of a policy has been estimated, the policy can be improved.

Markov decision problem

Policy improvement

- ▶ Once the value of a policy has been estimated, the policy can be improved.
- ▶ Let us assume we estimated V^π . Then, we compute:

$$\pi'(x) \leftarrow \arg \max_{a \in \mathcal{A}} \sum_{x'} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

for each state $x \in \mathcal{X}$.

Markov decision problem

Policy improvement

- ▶ Once the value of a policy has been estimated, the policy can be improved.
- ▶ Let us assume we estimated V^π . Then, we compute:

$$\pi'(x) \leftarrow \arg \max_{a \in \mathcal{A}} \sum_{x'} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^\pi(x')]$$

for each state $x \in \mathcal{X}$.

- ▶ Then π' is either better than π , or they are equivalent, that is:
 $V^{\pi'} \geq V^\pi$.

Markov decision problem



Example (continued)

Implement the policy improvement algorithm.

Apply it on the little MDP. How is the uniformly random policy improved?

Markov decision problem

Policy iteration

- ▶ Start with a random policy.
 - ▶ Then, alternate:
 - ▶ estimate the value of the current policy
 - ▶ improve the current policy
- until the value of the policy no longer improves.

This is the “policy iteration” algorithm [Howard, 1958].

Markov decision problem



Example (continued)

Implement policy iteration and test it on the little MDP and on the “21 with a dice” problem.

Markov decision problem

Value iteration

- ▶ The value V^* of the optimal policy π^* is:

$$V^*(x) \stackrel{\text{def}}{=} \max_{\pi} V^{\pi}(x), \forall x \in \mathcal{X}$$

Markov decision problem

Value iteration

- ▶ The value V^* of the optimal policy π^* is:

$$V^*(x) \stackrel{\text{def}}{=} \max_{\pi} V^{\pi}(x), \forall x \in \mathcal{X}$$

- ▶ It follows:

$$V^*(x) = \max_{a \in \mathcal{A}} \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^*(x')]$$

The “Bellman optimality equation”.

Markov decision problem

Value iteration

- ▶ The value V^* of the optimal policy π^* is:

$$V^*(x) \stackrel{\text{def}}{=} \max_{\pi} V^{\pi}(x), \forall x \in \mathcal{X}$$

- ▶ It follows:

$$V^*(x) = \max_{a \in \mathcal{A}} \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^*(x')]$$

The “Bellman optimality equation”.

- ▶ and an algorithm to obtain V^* directly:

- ▶ $k \leftarrow 0$
- ▶ initialize V_k
 - ▶ compute $V_{k+1} \leftarrow \max_{a \in \mathcal{A}} \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma V^*(x')]$ for all x
 - ▶ $k \leftarrow k + 1$
- ▶ repeat
- ▶ until $\|V_k - V_{k-1}\|_{\infty} \leq \frac{\epsilon(1-\gamma)}{2\gamma}$

named “value iteration”.

Markov decision problem



Example (continued)

Implement value iteration and test it on the little MDP and on the “21 with a dice” problem.

Markov decision problem

As a linear program

We can frame a Markov decision problem as a linear program:

$$\begin{aligned} & \min \sum_x V[x] \\ \text{s.t. } & V[x] - \sum_{x'} \mathcal{P}(x, a, x') (\mathcal{R}(x, a, x') + \gamma V[x']) \geq 0 \quad \forall (x, a) \in \mathcal{X} \times \mathcal{A} \end{aligned}$$

There is one constraint for each (state, action) pair and one variable per state.

The dual is:

$$\begin{aligned} & \max \sum_{(x, a)} \sum_{x'} \mathcal{P}(x, a, x') \mathcal{R}(x, a, x') \xi(x, a) \\ \text{s.t. } & \sum_a \xi(x', a) - \sum_{x, a} \gamma \mathcal{P}(x, a, x') \xi(x, a) \leq 1, \forall x' \in \mathcal{X} \\ & \text{and } \xi(x, a) \geq 0, \forall (x, a) \in \mathcal{X} \times \mathcal{A} \end{aligned}$$

Once solved, $\xi(x, a)$ is non zero if action a is optimal in x .

Very interesting theoretically speaking.

In practice, policy iteration is (usually) the best way to go.

Markov decision problem

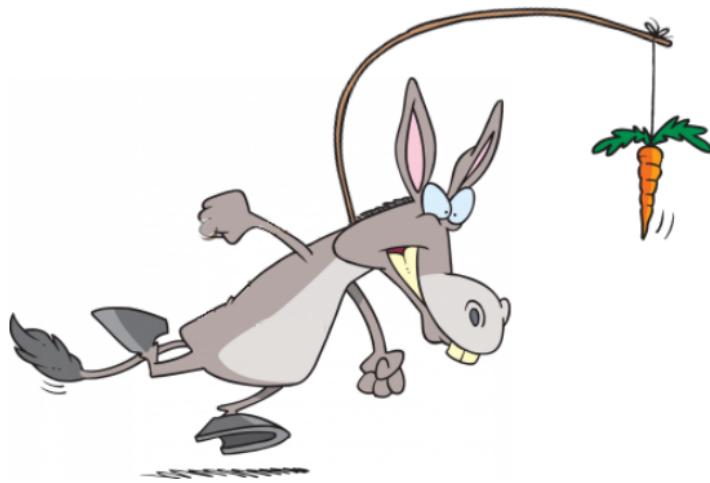


Example (continued)

Implement the LP approach on the little MDP and on the “21 with a dice” problem.

Hint: in Python, use the PULP package to solve an LP.

From MDPs to Reinforcement Learning



From MDPs to Reinforcement Learning

What if \mathcal{P} and \mathcal{R} are unknown?

From MDPs to Reinforcement Learning

The quality function

- ▶ Reminder: $V^\pi \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_{t>0} \sim \pi]$

Let us define $Q^\pi(x, a) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_0 = a, a_{t>0} \sim \pi]$

named the "**quality**" of the (state, action) pair (x, a) .

From MDPs to Reinforcement Learning

The quality function

- ▶ Reminder: $V^\pi \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_{t>0} \sim \pi]$

Let us define $Q^\pi(x, a) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x) | x_0 = x, a_0 = a, a_{t>0} \sim \pi]$

named the "**quality**" of the (state, action) pair (x, a) .

- ▶ We have: $V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) Q^\pi(x, a)$.

From MDPs to Reinforcement Learning

The quality function

► Reminder: $V^\pi \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x)|x_0 = x, a_{t>0} \sim \pi]$

Let us define $Q^\pi(x, a) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x)|x_0 = x, a_0 = a, a_{t>0} \sim \pi]$

named the "**quality**" of the (state, action) pair (x, a) .

► We have: $V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) Q^\pi(x, a)$.

► Bellman equation for Q :

$$Q^\pi(x, a) = \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma \sum_{a' \in \mathcal{A}} \pi(x', a') Q^\pi(x', a')].$$

From MDPs to Reinforcement Learning

The quality function

- ▶ Reminder: $V^\pi \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x)|x_0 = x, a_{t>0} \sim \pi]$

Let us define $Q^\pi(x, a) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x)|x_0 = x, a_0 = a, a_{t>0} \sim \pi]$

named the "**quality**" of the (state, action) pair (x, a) .

- ▶ We have: $V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) Q^\pi(x, a)$.
- ▶ Bellman equation for Q :
$$Q^\pi(x, a) = \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma \sum_{a' \in \mathcal{A}} \pi(x', a') Q^\pi(x', a')].$$
- ▶ As for V , we may define the optimal quality that is the quality of the optimal policy: $Q^* \stackrel{\text{def}}{=} \max_\pi Q^\pi$.

From MDPs to Reinforcement Learning

The quality function

- ▶ Reminder: $V^\pi \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x)|x_0 = x, a_{t>0} \sim \pi]$

Let us define $Q^\pi(x, a) \stackrel{\text{def}}{=} \mathbb{E}[\zeta(x)|x_0 = x, a_0 = a, a_{t>0} \sim \pi]$

named the "**quality**" of the (state, action) pair (x, a) .

- ▶ We have: $V^\pi(x) = \sum_{a \in \mathcal{A}} \pi(x, a) Q^\pi(x, a)$.

- ▶ Bellman equation for Q :

$$Q^\pi(x, a) =$$

$$\sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma \sum_{a' \in \mathcal{A}} \pi(x', a') Q^\pi(x', a')].$$

- ▶ As for V , we may define the optimal quality that is the quality of the optimal policy: $Q^* \stackrel{\text{def}}{=} \max_\pi Q^\pi$.

- ▶ Bellman optimality equation for Q^* :

$$Q^*(x, a) = \sum_{x' \in \mathcal{X}} \mathcal{P}(x, a, x') [\mathcal{R}(x, a, x') + \gamma \max_{a' \in \mathcal{A}} Q^*(x', a')].$$

Reinforcement Learning

Bellman equation and the TD error [Sutton, 1988]

Bellman approach:

- ▶ $\zeta = \sum_{t \geq 0} \gamma^t r_t, \gamma \in [0, 1[$

Reinforcement Learning

Bellman equation and the TD error [Sutton, 1988]

Bellman approach:

- ▶ $\zeta = \sum_{t \geq 0} \gamma^t r_t, \gamma \in [0, 1[$
- ▶ the value of a state is: $V^\pi(x) \stackrel{\text{def}}{=} \mathbb{E}_{a \sim \pi}(\zeta(x))$

Reminder: This quantifies what will happen to the agent in its future if it behaves according to π .

Reinforcement Learning

Bellman equation and the TD error [Sutton, 1988]

Bellman approach:

- ▶ $\zeta = \sum_{t \geq 0} \gamma^t r_t, \gamma \in [0, 1[$
- ▶ the value of a state is: $V^\pi(x) \stackrel{\text{def}}{=} \mathbb{E}_{a \sim \pi}(\zeta(x))$

Reminder: This quantifies what will happen to the agent in its future if it behaves according to π .

- ▶ re-written as: $V^\pi(x_t) = \mathbb{E}(r_t) + \gamma \mathbb{E}(V^\pi(x_{t+1}))$
sum of what will happen immediately + $\gamma \times$ what will happen then.

Reinforcement Learning

Bellman equation and the TD error [Sutton, 1988]

Bellman approach:

- ▶ $\zeta = \sum_{t \geq 0} \gamma^t r_t, \gamma \in [0, 1[$
- ▶ the value of a state is: $V^\pi(x) \stackrel{\text{def}}{=} \mathbb{E}_{a \sim \pi}(\zeta(x))$

Reminder: This quantifies what will happen to the agent in its future if it behaves according to π .

- ▶ re-written as: $V^\pi(x_t) = \mathbb{E}(r_t) + \gamma \mathbb{E}(V^\pi(x_{t+1}))$

sum of what will happen immediately + $\gamma \times$ what will happen then.

- ▶ at time t , $r_t + \gamma(V^\pi(x_{t+1})) - V^\pi(x_t)$

is an estimation of the error of estimation of V : the difference between what we expected before performing a_t ($= V^\pi(x_t)$) and what we can expect now that we performed a_t and observed the x_{t+1} and r_t ($= r_t + \gamma(V^\pi(x_{t+1}))$). TD-error

This TD-error may be used to learn the optimal behavior.

Reinforcement Learning

The temporal difference

- ▶ computing V by gradient descent:

$$V(x_{t+1}) \leftarrow V(x_t) - \eta[r_t + \gamma(V_t(x_{t+1})) - V(x_t)]$$

where η is a learning rate, adequately decreasing along time.

Reinforcement Learning

The temporal difference

- ▶ computing V by gradient descent:

$$V(x_{t+1}) \leftarrow V(x_t) - \eta[r_t + \gamma(V_t(x_{t+1})) - V(x_t)]$$

where η is a learning rate, adequately decreasing along time.

- ▶ We may also consider the quality of an (x, a) pair:

$$Q(x_t, a_t) \stackrel{\text{def}}{=} \mathbb{E}(\zeta(x_t | a_t = a, \pi))$$

Reinforcement Learning

The temporal difference

- ▶ computing V by gradient descent:

$$V(x_{t+1}) \leftarrow V(x_t) - \eta[r_t + \gamma(V_t(x_{t+1})) - V(x_t)]$$

where η is a learning rate, adequately decreasing along time.

- ▶ We may also consider the quality of an (x, a) pair:

$$Q(x_t, a_t) \stackrel{\text{def}}{=} \mathbb{E}(\zeta(x_t | a_t = a, \pi))$$

- ▶ At t , we may consider $r_t + \gamma \max_{a'} Q(x_{t+1}, a') - Q(x_t, a_t)$ as a correction to $Q(x_t, a_t)$.

\rightsquigarrow

$$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \eta[r_t + \gamma \max_b Q(x_{t+1}, b) - Q(x_t, a_t)]$$

Reinforcement Learning

The temporal difference

- ▶ computing V by gradient descent:

$$V(x_{t+1}) \leftarrow V(x_t) - \eta[r_t + \gamma(V_t(x_{t+1})) - V(x_t)]$$

where η is a learning rate, adequately decreasing along time.

- ▶ We may also consider the quality of an (x, a) pair:

$$Q(x_t, a_t) \stackrel{\text{def}}{=} \mathbb{E}(\zeta(x_t | a_t = a, \pi))$$

- ▶ At t , we may consider $r_t + \gamma \max_{a'} Q(x_{t+1}, a') - Q(x_t, a_t)$ as a correction to $Q(x_t, a_t)$.

\rightsquigarrow

$$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \eta[r_t + \gamma \max_b Q(x_{t+1}, b) - Q(x_t, a_t)]$$

- ▶ \rightsquigarrow learning π^* algorithm.

Reinforcement Learning

Sketch of an RL algorithm to learn π^*

The goal of this algorithm is to estimate Q^* by interacting with the environment and correcting its estimate of Q^* .

1. initialize Q .
2. set the agent in some random initial state $\in \mathcal{X}_0$.
3. run the agent in the environment:
at each step, record the state x_t , the action performed a_t , the reward collected r_t , and the next state x_{t+1} .
4. correct Q .
5. continue until a terminal state is reached.
6. do it again and again (re-starting at step 2).

Reinforcement Learning

Q-Learning [Watkins, 1989]

$Q(x, a) \leftarrow$ some value (0, random, ...)

repeat

$t \leftarrow 0$

Initialize the state of the agent x_t

while episode not completed **do**

choose an action to perform in state x_t : a_t

perform this action and observe r_t and x_{t+1}

update:

$Q(x_t, a_t) \leftarrow Q(x_t, a_t) + \alpha[r_t + \max_{b \in \mathcal{A}} Q(x_{t+1}, b) - Q(x_t, a_t)]$

$t++$

end while

until some condition is met.

At completion (if you looped enough): $\pi^*(x) = \arg \max_a Q(x, a), \forall x$

Reinforcement Learning

About Q-Learning

At completion (if you looped enough): $\pi^*(x) = \arg \max_a Q(x, a), \forall x$

- ▶ Remark: you never know if you looped enough!
- ▶ Asymptotic convergence to Q^* .

- ▶ α depends on x_t and a_t .
- ▶ for each (x, a) , we should have:
$$\sum_{\{t \text{ at which } (x, a) \text{ is visited}\}} \alpha_t(s, n(s)) = +\infty$$
 and
$$\sum_{\{t \text{ at which } (x, a) \text{ is visited}\}} \alpha_t^2(s, n(s)) < +\infty$$
e.g. $\alpha(x, a) \equiv \frac{1}{n(x, a) + 1}$

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”
How do we do that?

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”
How do we do that?
- ▶ No perfect solution.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”
How do we do that?
- ▶ No perfect solution.
- ▶ Intuition:

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”
How do we do that?
- ▶ No perfect solution.
- ▶ Intuition:
 - ▶ initially, we do not know anything about Q^* . We have to test the effect of the actions: we have to **explore**.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”
How do we do that?
- ▶ No perfect solution.
- ▶ Intuition:
 - ▶ initially, we do not know anything about Q^* . We have to test the effect of the actions: we have to explore.
 - ▶ step by step, we acquire some knowledge about which actions are good, and which are bad. We can **exploit** this knowledge.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”
How do we do that?
- ▶ No perfect solution.
- ▶ Intuition:
 - ▶ initially, we do not know anything about Q^* . We have to test the effect of the actions: we have to explore.
Select action at random
 - ▶ step by step, we acquire some knowledge about which actions are good, and which are bad. We can exploit this knowledge.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”
How do we do that?
- ▶ No perfect solution.
- ▶ Intuition:
 - ▶ initially, we do not know anything about Q^* . We have to test the effect of the actions: we have to explore.
Select action at random
 - ▶ step by step, we acquire some knowledge about which actions are good, and which are bad. We can exploit this knowledge.
Select the “best” action: $\arg \max_a Q(x_t, a)$.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”
How do we do that?
- ▶ No perfect solution.
- ▶ Intuition:
 - ▶ initially, we do not know anything about Q^* . We have to test the effect of the actions: we have to explore.
Select action at random
 - ▶ step by step, we acquire some knowledge about which actions are good, and which are bad. We can exploit this knowledge.
Select the “best” action: $\arg \max_a Q(x_t, a)$.
- ▶ Exploration and exploitation should be carefully balanced.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:
 - ▶ $\epsilon \leftarrow 1$.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:
 - ▶ $\epsilon \leftarrow 1$.
 - ▶ Select action $\arg \max_a Q(x_t, a)$ with probability ϵ and a random action otherwise.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:
 - ▶ $\epsilon \leftarrow 1$.
 - ▶ Select action $\arg \max_a Q(x_t, a)$ with probability ϵ and a random action otherwise.
 - ▶ Slowly decrease ϵ along the episodes (e.g. $\epsilon \leftarrow 0.98\epsilon$).

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:
 - ▶ $\epsilon \leftarrow 1$.
 - ▶ Select action $\arg \max_a Q(x_t, a)$ with probability ϵ and a random action otherwise.
 - ▶ Slowly decrease ϵ along the episodes (e.g. $\epsilon \leftarrow 0.98\epsilon$).
- ▶ The softmax strategy:

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:
 - ▶ $\epsilon \leftarrow 1$.
 - ▶ Select action $\arg \max_a Q(x_t, a)$ with probability ϵ and a random action otherwise.
 - ▶ Slowly decrease ϵ along the episodes (e.g. $\epsilon \leftarrow 0.98\epsilon$).
- ▶ The softmax strategy:
 - ▶ $\tau \leftarrow 1000$.

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:
 - ▶ $\epsilon \leftarrow 1$.
 - ▶ Select action $\arg \max_a Q(x_t, a)$ with probability ϵ and a random action otherwise.
 - ▶ Slowly decrease ϵ along the episodes (e.g. $\epsilon \leftarrow 0.98\epsilon$).
- ▶ The softmax strategy:
 - ▶ $\tau \leftarrow 1000$.
 - ▶ Each time you need to select an action:

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:
 - ▶ $\epsilon \leftarrow 1$.
 - ▶ Select action $\arg \max_a Q(x_t, a)$ with probability ϵ and a random action otherwise.
 - ▶ Slowly decrease ϵ along the episodes (e.g. $\epsilon \leftarrow 0.98\epsilon$).
- ▶ The softmax strategy:
 - ▶ $\tau \leftarrow 1000$.
 - ▶ Each time you need to select an action:
 - ▶ compute $p(a) \leftarrow \frac{e^{\frac{Q(x_t, a)}{\tau}}}{\sum_{a' \in \mathcal{A}} e^{\frac{Q(x_t, a')}{\tau}}}$,

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:
 - ▶ $\epsilon \leftarrow 1$.
 - ▶ Select action $\arg \max_a Q(x_t, a)$ with probability ϵ and a random action otherwise.
 - ▶ Slowly decrease ϵ along the episodes (e.g. $\epsilon \leftarrow 0.98\epsilon$).
- ▶ The softmax strategy:
 - ▶ $\tau \leftarrow 1000$.
 - ▶ Each time you need to select an action:
 - ▶ compute $p(a) \leftarrow \frac{e^{\frac{Q(x_t, a)}{\tau}}}{\sum_{a' \in \mathcal{A}} e^{\frac{Q(x_t, a')}{\tau}}}$,
 - ▶ draw an action at random according to p .

Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:
 - ▶ $\epsilon \leftarrow 1$.
 - ▶ Select action $\arg \max_a Q(x_t, a)$ with probability ϵ and a random action otherwise.
 - ▶ Slowly decrease ϵ along the episodes (e.g. $\epsilon \leftarrow 0.98\epsilon$).
- ▶ The softmax strategy:
 - ▶ $\tau \leftarrow 1000$.
 - ▶ Each time you need to select an action:
 - ▶ compute $p(a) \leftarrow \frac{e^{\frac{Q(x_t, a)}{\tau}}}{\sum_{a' \in \mathcal{A}} e^{\frac{Q(x_t, a')}{\tau}}}$,
 - ▶ draw an action at random according to p .
 - ▶ Slowly decrease τ along the episodes.

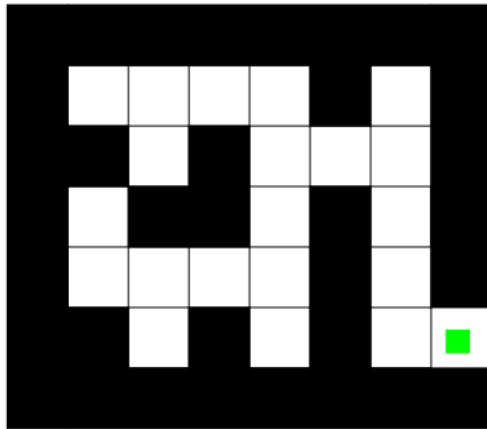
Reinforcement Learning

About Q-Learning

- ▶ “choose an action to perform in state x_t : a_t ”.
- ▶ explore and progressively exploit more and more.
- ▶ The ϵ -decreasing greedy strategy:
 - ▶ $\epsilon \leftarrow 1$.
 - ▶ Select action $\arg \max_a Q(x_t, a)$ with probability ϵ and a random action otherwise.
 - ▶ Slowly decrease ϵ along the episodes (e.g. $\epsilon \leftarrow 0.98\epsilon$).
- ▶ The softmax strategy:
 - ▶ $\tau \leftarrow 1000$.
 - ▶ Each time you need to select an action:
 - ▶ compute $p(a) \leftarrow \frac{e^{\frac{Q(x_t, a)}{\tau}}}{\sum_{a' \in \mathcal{A}} e^{\frac{Q(x_t, a')}{\tau}}}$,
 - ▶ draw an action at random according to p .
 - ▶ Slowly decrease τ along the episodes.
 - ▶ Rationale: when τ is large, uniformly random selection. τ close to 0, greedy selection.

Reinforcement Learning

Q-Learning in action: escaping a maze



Goal: reach the green cell from any location.

Question 1: propose a Markov decision process: what are $(\mathcal{T}, \mathcal{X}, \mathcal{X}_0, \mathcal{X}_f, \mathcal{A}, \mathcal{P}, \mathcal{R})$?

Question 2: model this task as a Markov decision problem: what is ζ ?

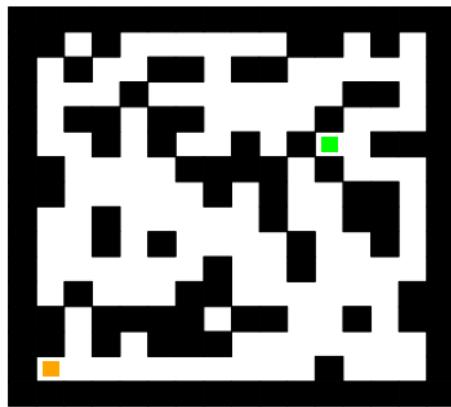
Reinforcement Learning

Q-Learning in action: escaping a maze

We use an extremely basic Q-Learning.

Has a very local perception: sees only the 4 neighboring cells.

Gets $r = 100$ when reaching the goal, 0 otherwise.



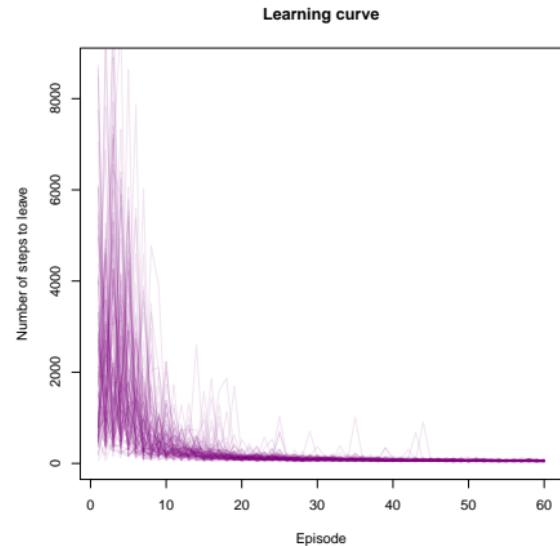
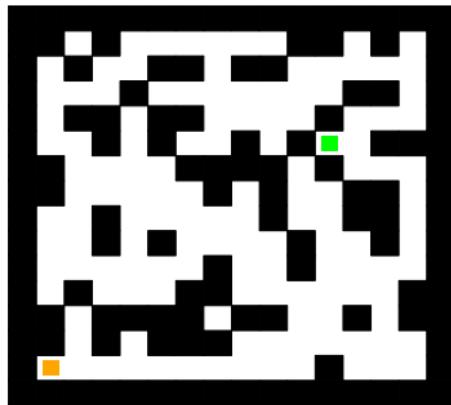
Reinforcement Learning

Q-Learning in action: escaping a maze

We use an extremely basic Q-Learning.

Has a very local perception: sees only the 4 neighboring cells.

Gets $r = 100$ when reaching the goal, 0 otherwise.



Reinforcement Learning

Q-Learning in action

1st reach



Reinforcement Learning

Q-Learning in action

1st reach



10th reach



Reinforcement Learning

Q-Learning in action

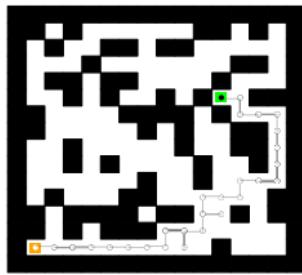
1st reach



10th reach



60th reach



Reinforcement Learning

Not only the last step matters: Eligibility traces

- ▶ When reaching the goal for the first time, only the last (x, a) pair is modified = rewarded.
- ▶ Previous pairs also had a role since they led to the last pair.
- ▶ They may be rewarded too.
- ▶ Less and less as they are further away the last pair.

¹very weird name.

Reinforcement Learning

Not only the last step matters: Eligibility traces

- ▶ When reaching the goal for the first time, only the last (x, a) pair is modified = rewarded.
- ▶ Previous pairs also had a role since they led to the last pair.
- ▶ They may be rewarded too.
- ▶ Less and less as they are further away the last pair.

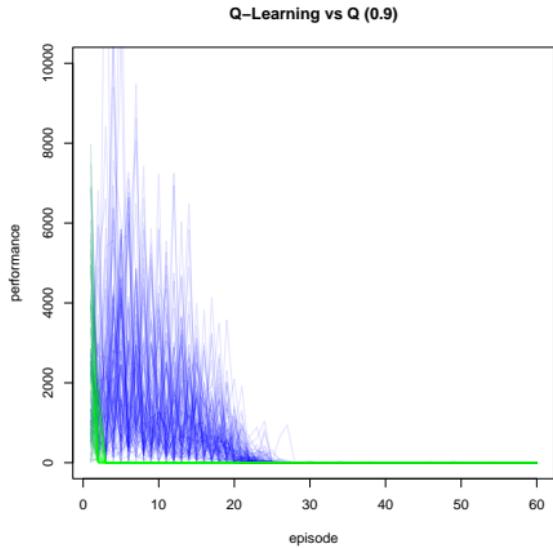
→ Eligibility traces¹:

- ▶ each (x, a) pair has an associated eligibility $e(x, a)$.
- ▶ $Q(x, a)$ is updated for each (x, a) pair that is visited.
- ▶ $e(x, a)$ is set to 1 when this pair is visited.
- ▶ Then, $e(x, a)$ is decaying along time: $e(x, a) \leftarrow \lambda e(x, a)$.

¹very weird name.

Reinforcement Learning

Not only the last step matters: Eligibility traces



100 runs of Q-Learning vs. 100 runs of Q ($\lambda = 0.9$).

Q (λ)

initialize Q

repeat

$$e(x, a) \leftarrow 0, \forall (x, a) \in \mathcal{X} \times \mathcal{A}$$

$$t \leftarrow 0$$

initialize x_0 and choose action a_0

repeat

emit a_t and observe r_t and x_{t+1}

choose action a_{t+1} to be emitted in x_{t+1}

$$a_{t+1} \leftarrow \arg \max_{a \in \mathcal{A}} Q(x_{t+1}, a)$$

$$\delta \leftarrow r_t + \gamma \hat{Q}(x_{t+1}, a_{t+1}) - Q(x_t, a_t)$$

$$e(x_t, a_t) \leftarrow e(x_t, a_t) + 1$$

for $(x, a) \in \mathcal{X} \times \mathcal{A}$ **do**

$$Q(x, a) \leftarrow Q(x, a) + \alpha \delta e(x, a)$$

$$e(x, a) \leftarrow \gamma \lambda e(x, a)$$

end for

$$t \leftarrow t + 1$$

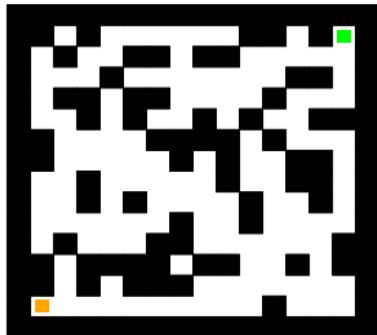
until end of episode

until ...

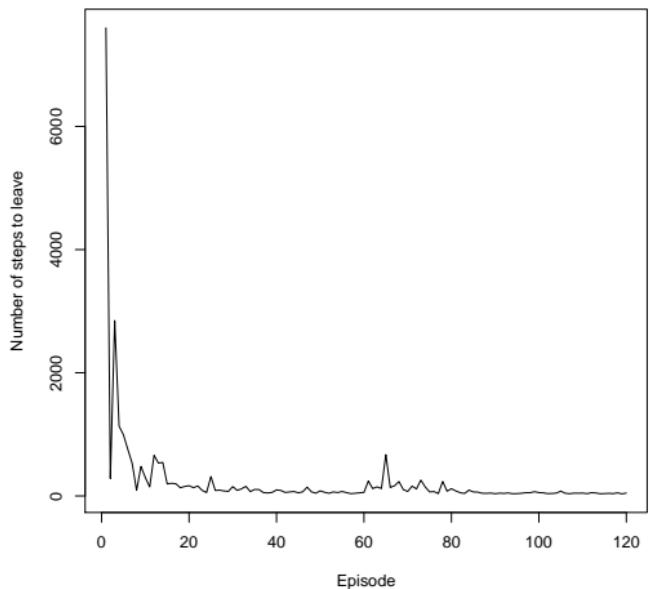
Reinforcement Learning

Q-Learning continuously adapts to its environment

The goal state moves nearby:



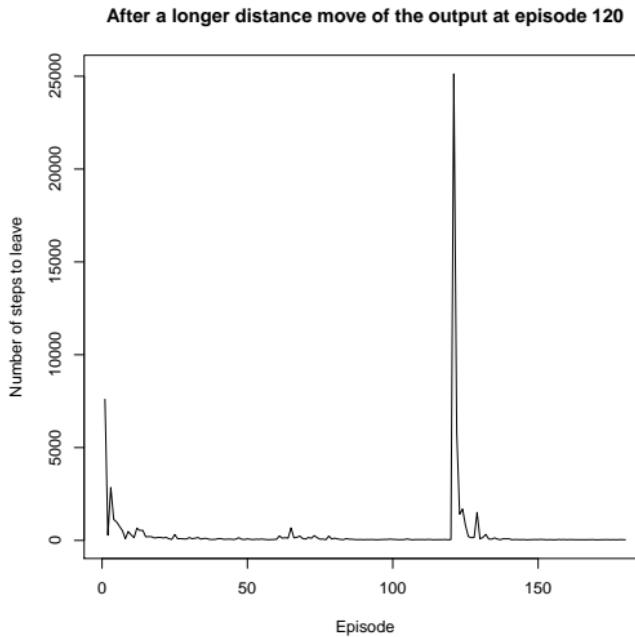
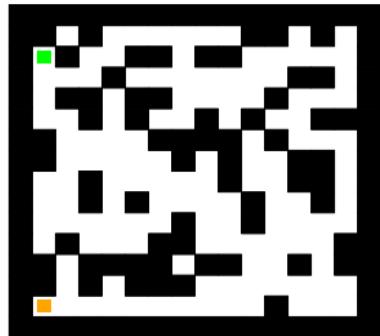
After a small distance move of the output at episode 60



Reinforcement Learning

Q-Learning continuously adapts to its environment

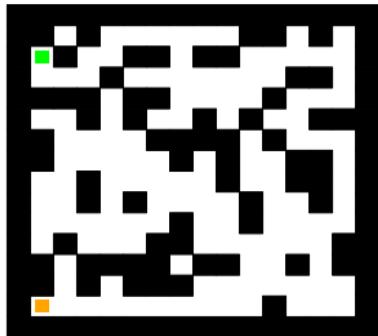
The goal state moves farther away:



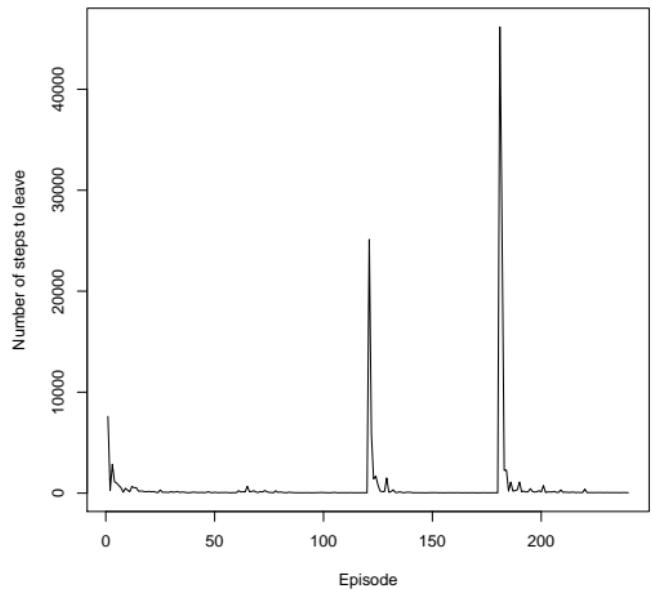
Reinforcement Learning

Q-Learning continuously adapts to its environment

Blocking the path:



After adding a wall on the path at episode 180



Reinforcement Learning

From table to function approximation

- ▶ This is the “tabular” Q-Learning: Q is represented in a “table”.

Reinforcement Learning

From table to function approximation

- ▶ This is the “tabular” Q-Learning: Q is represented in a “table”.
- ▶ What about large \mathcal{X} ?

Reinforcement Learning

From table to function approximation

- ▶ This is the “tabular” Q-Learning: Q is represented in a “table”.
- ▶ What about large \mathcal{X} ?
- ▶ Impossible to store Q in a table.

Reinforcement Learning

From table to function approximation

- ▶ This is the “tabular” Q-Learning: Q is represented in a “table”.
- ▶ What about large \mathcal{X} ?
- ▶ Impossible to store Q in a table.
- ▶ Use a function approximator, that is, replace the table $Q[x, a]$ by a function $Q(x, a)$.

Reinforcement Learning

From table to function approximation

- ▶ This is the “tabular” Q-Learning: Q is represented in a “table”.
- ▶ What about large \mathcal{X} ?
- ▶ Impossible to store Q in a table.
- ▶ Use a function approximator, that is, replace the table Q [x , a] by a function $Q(x, a)$.
- ▶ $Q(x, a)$ returns an estimate of $Q(x, a)$.

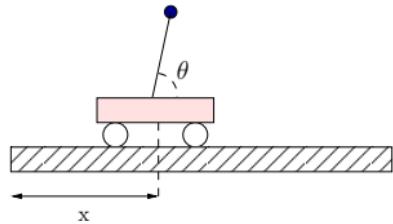
Reinforcement Learning

From table to function approximation

- ▶ This is the “tabular” Q-Learning: Q is represented in a “table”.
- ▶ What about large \mathcal{X} ?
- ▶ Impossible to store Q in a table.
- ▶ Use a function approximator, that is, replace the table Q [x , a] by a function $Q(x, a)$.
- ▶ $Q(x, a)$ returns an estimate of $Q(x, a)$.
- ▶ This estimate may be updated/improved by learning.

Reinforcement Learning

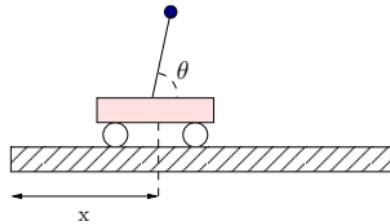
Value function



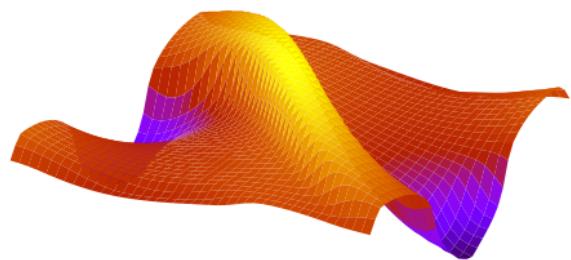
State is $(\theta, \dot{\theta})$
Action is $\ddot{\theta}$

Reinforcement Learning

Value function



State is $(\theta, \dot{\theta})$
Action is $\ddot{\theta}$



$(\theta, \dot{\theta})$ plane
 z is $V(x)$
Maximize value \rightsquigarrow reach the top of V

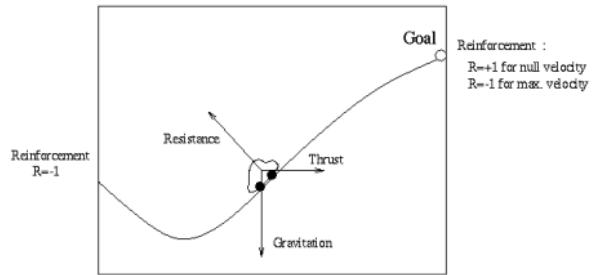
Reinforcement Learning

Handling large \mathcal{X} : the function approximator zoo

- ▶ neural network [Lin, 1991; Riedmiller, 2005; ...],
- ▶ random forest [Geurts *et al.*, 2006],
- ▶ SVM and kernels,
- ▶ and many other ideas from statistical learning (supervised learning).
- ▶ Tabular with progressive and adaptive state partitioning.

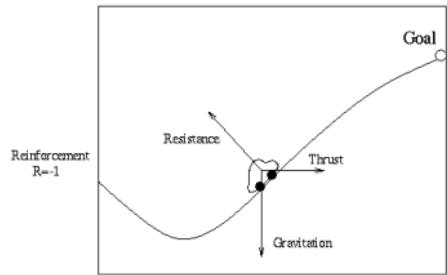
Reinforcement Learning

Progressive and adaptive state partitioning [Munos, Moore, MLJ, 2001]

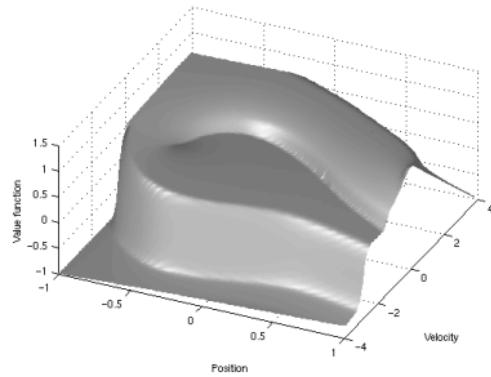


Reinforcement Learning

Progressive and adaptive state partitioning [Munos, Moore, MLJ, 2001]

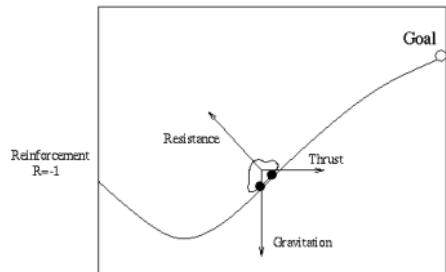


Reinforcement :
 $R=+1$ for null velocity
 $R=-1$ for max. velocity

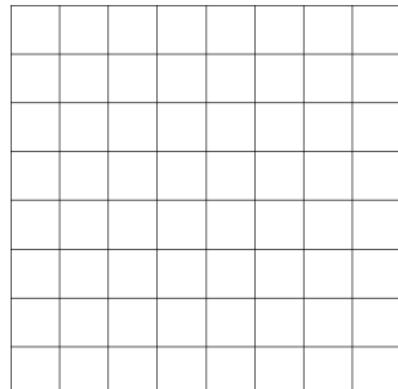
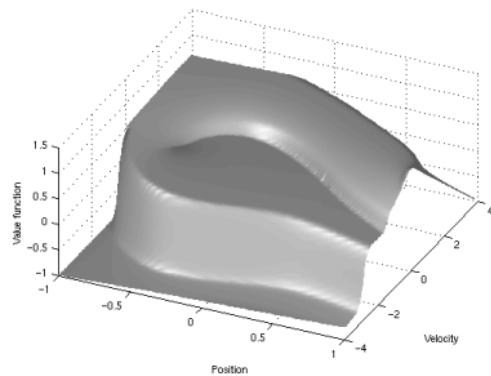


Reinforcement Learning

Progressive and adaptive state partitioning [Munos, Moore, MLJ, 2001]

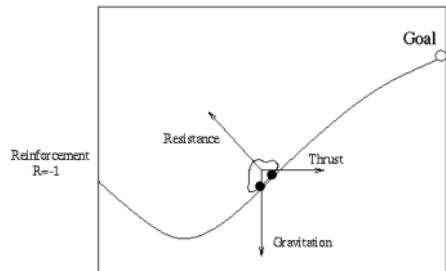


Reinforcement :
 $R=+1$ for null velocity
 $R=-1$ for max. velocity

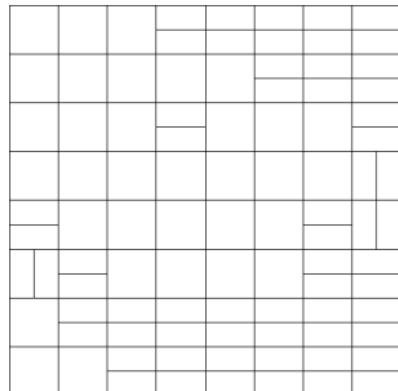
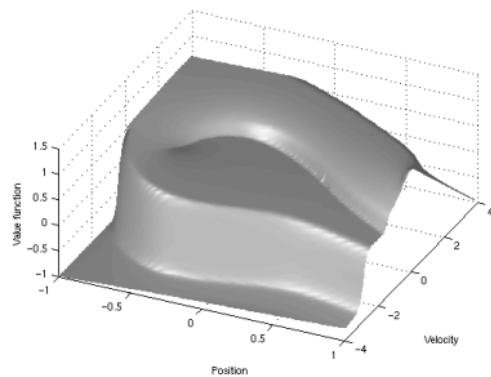


Reinforcement Learning

Progressive and adaptive state partitioning [Munos, Moore, MLJ, 2001]

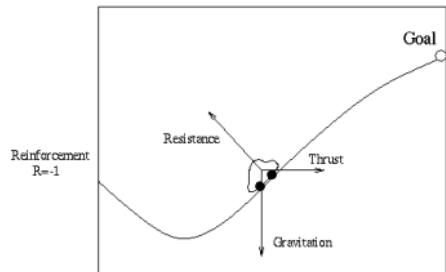


Reinforcement :
 $R=+1$ for null velocity
 $R=-1$ for max. velocity

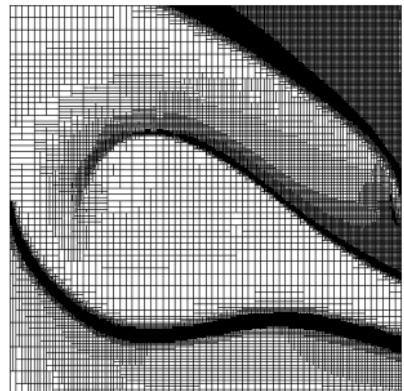
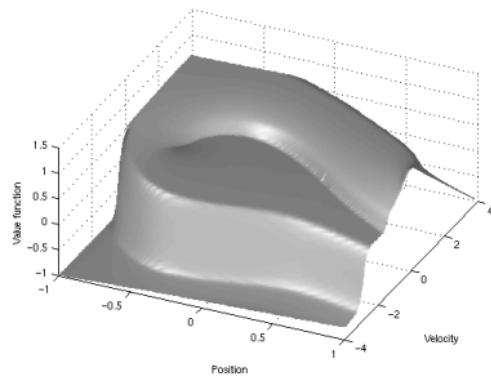


Reinforcement Learning

Progressive and adaptive state partitioning [Munos, Moore, MLJ, 2001]



Reinforcement :
 $R=+1$ for null velocity
 $R=-1$ for max. velocity



Reinforcement Learning

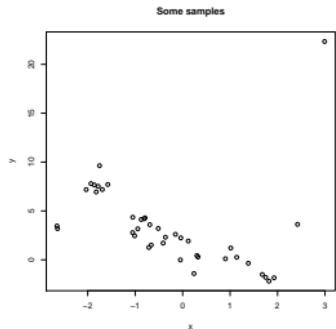
Neural Q-Learning

Trendy people call that "deep RL" though there is only shallow networks most of the times.

- ▶ Use of a neural network to represent Q .
- ▶ Input = the current state
- ▶ Output = either an estimate of $Q(x, a), \forall a$, or an estimate of $\pi(x)$.

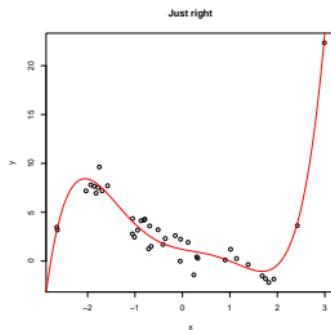
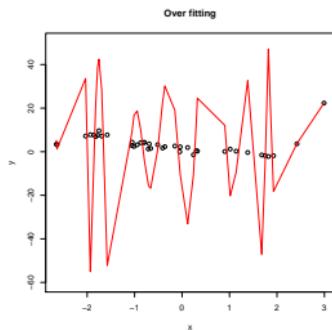
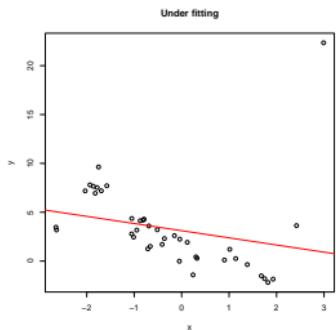
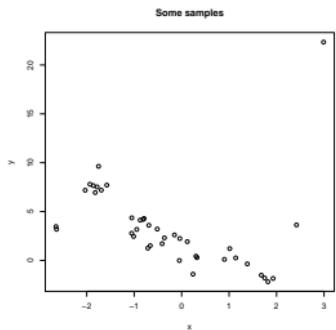
Reminder: function approximation

Reminder: function approximation

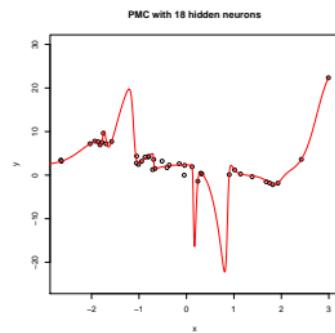
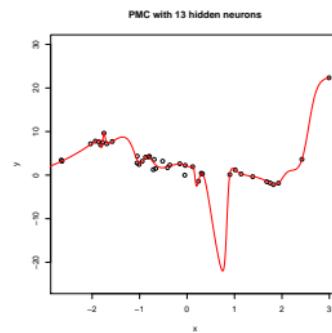
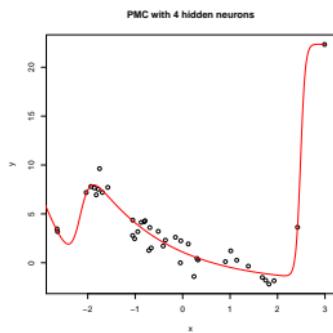
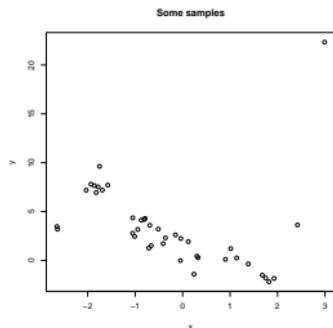


A few samples.

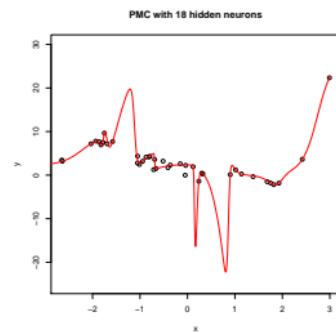
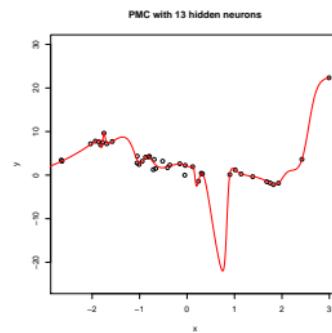
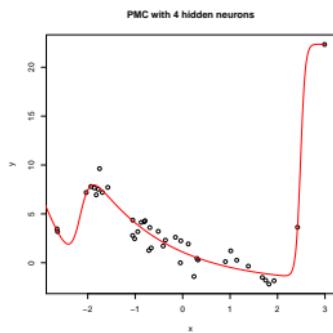
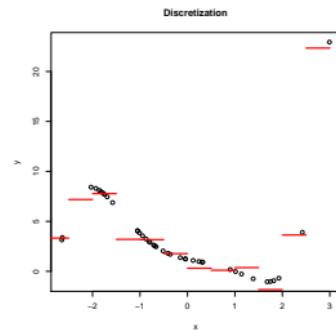
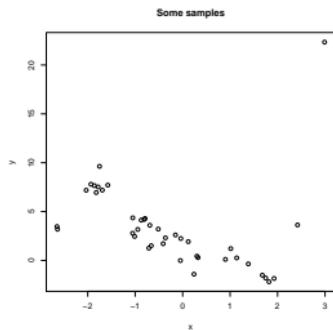
Reminder: function approximation



Reminder: function approximation



Reminder: function approximation



Reinforcement Learning

Neural Q-Learning

Tabular Q-Learning

```
Initialize the Q-table  $Q$ .  
repeat  
    set initial state  $x_0$   
     $t \leftarrow 0$   
    while episode not complete  
        do  
            • select  $a_t$  and emit it  
            • observe  $r_t$  and  $x_{t+1}$   
            • update  $Q(x_t, a_t)$   
            •  $t \leftarrow t + 1$   
        end while  
until ...
```

DQN

```
Initialize parameters  $\theta$   $Q$   
repeat  
    set initial state  $x_0$   
     $t \leftarrow 0$   
    while episode not complete  
        do  
            • select  $a_t$  and emit it  
            • observe  $r_t$  and  $x_{t+1}$   
            •  
        update  $\theta$  to minimise the prediction error (temporal difference) for  
         $(x_t, a_t)$   
        •  $t \leftarrow t + 1$   
    end while  
until ...
```

Reinforcement Learning

Neural Q-Learning

There are many things that can go wrong with this algorithm.

May be painful to debug.

- ▶ Design of the state space: the state is a vector; neighboring vectors should be neighbor states.

Reinforcement Learning

Neural Q-Learning

There are many things that can go wrong with this algorithm.

May be painful to debug.

- ▶ Design of the state space: the state is a vector; neighboring vectors should be neighbor states.
- ▶ This algorithm is very unstable: every little thing can make it fail, or perform poorly.

Reinforcement Learning

Neural Q-Learning

There are many things that can go wrong with this algorithm.

May be painful to debug.

- ▶ Design of the state space: the state is a vector; neighboring vectors should be neighbor states.
- ▶ This algorithm is very unstable: every little thing can make it fail, or perform poorly.
- ▶ Today, nobody uses this naive online version.

Reinforcement Learning

Neural Q-Learning

There are many things that can go wrong with this algorithm.

May be painful to debug.

- ▶ Design of the state space: the state is a vector; neighboring vectors should be neighbor states.
- ▶ This algorithm is very unstable: every little thing can make it fail, or perform poorly.
- ▶ Today, nobody uses this naive online version.
- ▶ Many tricks have been introduced recently:

Reinforcement Learning

Neural Q-Learning

There are many things that can go wrong with this algorithm.

May be painful to debug.

- ▶ Design of the state space: the state is a vector; neighboring vectors should be neighbor states.
- ▶ This algorithm is very unstable: every little thing can make it fail, or perform poorly.
- ▶ Today, nobody uses this naive online version.
- ▶ Many tricks have been introduced recently:
 - ▶ online → batch,

Reinforcement Learning

Neural Q-Learning

There are many things that can go wrong with this algorithm.

May be painful to debug.

- ▶ Design of the state space: the state is a vector; neighboring vectors should be neighbor states.
- ▶ This algorithm is very unstable: every little thing can make it fail, or perform poorly.
- ▶ Today, nobody uses this naive online version.
- ▶ Many tricks have been introduced recently:
 - ▶ online → batch,
 - ▶ use of a replay buffer,

Reinforcement Learning

Neural Q-Learning

There are many things that can go wrong with this algorithm.

May be painful to debug.

- ▶ Design of the state space: the state is a vector; neighboring vectors should be neighbor states.
- ▶ This algorithm is very unstable: every little thing can make it fail, or perform poorly.
- ▶ Today, nobody uses this naive online version.
- ▶ Many tricks have been introduced recently:
 - ▶ online → batch,
 - ▶ use of a replay buffer,
 - ▶ use of a target buffer,

Reinforcement Learning

Neural Q-Learning

There are many things that can go wrong with this algorithm.

May be painful to debug.

- ▶ Design of the state space: the state is a vector; neighboring vectors should be neighbor states.
- ▶ This algorithm is very unstable: every little thing can make it fail, or perform poorly.
- ▶ Today, nobody uses this naive online version.
- ▶ Many tricks have been introduced recently:
 - ▶ online → batch,
 - ▶ use of a replay buffer,
 - ▶ use of a target buffer,
 - ▶ prioritize samples in the replay buffer.

Reinforcement Learning

Neural Q-Learning

There are many things that can go wrong with this algorithm.

May be painful to debug.

- ▶ Design of the state space: the state is a vector; neighboring vectors should be neighbor states.
- ▶ This algorithm is very unstable: every little thing can make it fail, or perform poorly.
- ▶ Today, nobody uses this naive online version.
- ▶ Many tricks have been introduced recently:
 - ▶ online → batch,
 - ▶ use of a replay buffer,
 - ▶ use of a target buffer,
 - ▶ prioritize samples in the replay buffer.
 - ▶ and more.

Reinforcement Learning

Neural Q-Learning/DQN

There are many things that can go wrong with this algorithm.
May be painful to debug.

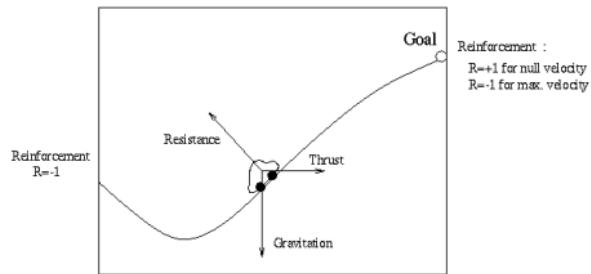
Neural networks may be tricky to train in supervised learning.
In RL, NN are used in a feedback loop: feedback loops are well-known to be very sensible, prone to very quickly (exponential rate) turn tiny errors into catastrophes.

When coding DQN:

- ▶ Each part of the algorithm should be tested.
- ▶ One should check that the NN is able to regress the Q function.

Reinforcement Learning

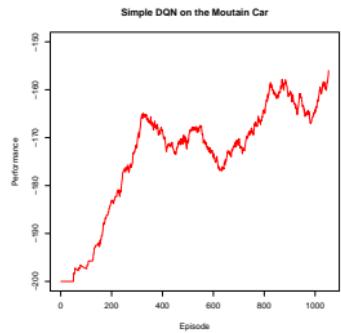
Neural Q-Learning/DQN



Reinforcement Learning

Neural Q-Learning/DQN

On the mountain-car:

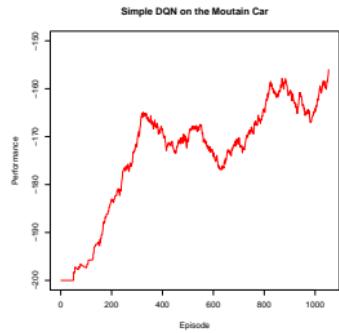


red: 128 neurons

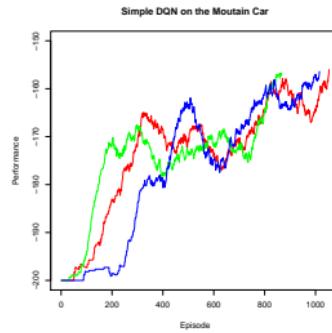
Reinforcement Learning

Neural Q-Learning/DQN

On the mountain-car:



red: 128 neurons

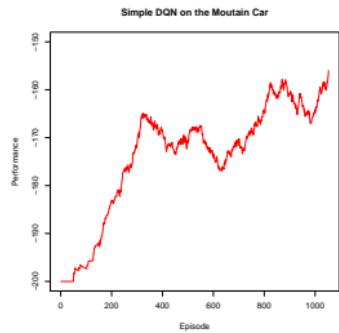


green: 32 neurons,
blue: 64 neurons

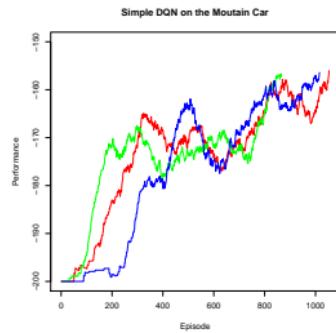
Reinforcement Learning

Neural Q-Learning/DQN

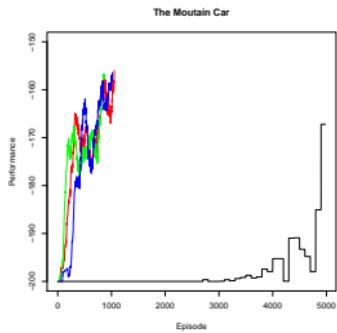
On the mountain-car:



red: 128 neurons



green: 32 neurons,
blue: 64 neurons



+ tabular QL

Reinforcement Learning

Neural Q-Learning/DQN: state = an image [Mnih et al., 2015]

The Atari game suite:



- ▶ Input: image
 - CNN
 - fully connected layers
 - P outputs, one per action.
- ▶ DQN learns to play 49 games.
- ▶ The CNN part learns an embedding of the state (image) that is used to learn the value of each action.

Reinforcement Learning

Policy gradient approach

An other family of RL algorithms

Reinforcement Learning

Policy gradient approach

- ▶ Do we have to learn a value function to learn a policy?

Reinforcement Learning

Policy gradient approach

- ▶ Do we have to learn a value function to learn a policy?
- ▶ No.

Reinforcement Learning

Policy gradient approach

- ▶ Do we have to learn a value function to learn a policy?
- ▶ No.
- ▶ → policy gradient.

Reinforcement Learning

Policy gradient approach

- ▶ Do we have to learn a value function to learn a policy?
- ▶ No.
- ▶ → policy gradient.
- ▶ Idea:

Reinforcement Learning

Policy gradient approach

- ▶ Do we have to learn a value function to learn a policy?
- ▶ No.
- ▶ → policy gradient.
- ▶ Idea:
 - ▶ Policy represented by a NN with parameters θ .

Reinforcement Learning

Policy gradient approach

- ▶ Do we have to learn a value function to learn a policy?
- ▶ No.
- ▶ → policy gradient.
- ▶ Idea:
 - ▶ Policy represented by a NN with parameters θ .
 - ▶ Run one or more trajectories with this policy.

Reinforcement Learning

Policy gradient approach

- ▶ Do we have to learn a value function to learn a policy?
- ▶ No.
- ▶ → policy gradient.
- ▶ Idea:
 - ▶ Policy represented by a NN with parameters θ .
 - ▶ Run one or more trajectories with this policy.
 - ▶ We can show that maximizing ζ is equivalent to minimizing $\mathcal{L}(\theta) \equiv -Q \log (\pi^\theta)$.

Reinforcement Learning

Policy gradient approach: REINFORCE [Williams, 1992]

initialize θ

repeat

 perform K episodes and collect transitions $(x_{k,t}, a_{k,t}, r_{k,t}, x_{k,t+1})$

for for each episode k **do**

for for each step t in episode k **do**

$$Q_{k,t} \leftarrow \sum_{j=t}^{j=|\tau_k|-1} \gamma^{j-t} r_{k,j}$$

$$\theta \leftarrow \theta + \alpha \gamma^t Q_{k,t} \nabla_\theta \log(\pi^\theta(s_{k,t}, a_{k,t}))$$

end for

end for

until ...

Reinforcement Learning

Policy gradient approach

REINFORCE suffers from important drawbacks:

- ▶ REINFORCE needs full episodes.
- ▶ REINFORCE is unstable because the update have a large variance.
- ▶ there is no explicit exploration.

Reinforcement Learning

Policy gradient approach

REINFORCE suffers from important drawbacks:

- ▶ REINFORCE needs full episodes.
However, we do not care about returns when γ^k vanishes.
- ▶ REINFORCE is unstable because the update have a large variance.
- ▶ there is no explicit exploration.

Reinforcement Learning

Policy gradient approach

REINFORCE suffers from important drawbacks:

- ▶ REINFORCE needs full episodes.
However, we do not care about returns when γ^k vanishes.
- ▶ REINFORCE is unstable because the update have a large variance.
Variance may be lowered by replacing $Q_{k,t}$ by $Q - b(x_{k,t})$.
- ▶ there is no explicit exploration.

Reinforcement Learning

Policy gradient approach

REINFORCE suffers from important drawbacks:

- ▶ REINFORCE needs full episodes.
However, we do not care about returns when γ^k vanishes.
- ▶ REINFORCE is unstable because the update have a large variance.
Variance may be lowered by replacing $Q_{k,t}$ by $Q - b(x_{k,t})$.
- ▶ there is no explicit exploration.
The trouble comes from the degenerescence of $\pi(x, .)$.
We can add an extra term in \mathcal{L} to force the non degenerescence of $\pi(x, .)$: maximize the entropy of $\pi(x, .) \rightarrow$
 $\mathcal{L}(\theta) \equiv -Q \log (\pi^\theta) + \lambda H(\pi^\theta)$.
Regularization is a key ingredient in optimization.

Reinforcement Learning

Policy search approach

One may use other types of optimization algorithm to optimize ζ .

Some researchers have investigated the use of evolutionary algorithms.

Reinforcement Learning

Actor-critic

The combination of the goods of value based and policy gradient approaches.

Reinforcement Learning

Actor-critic

The combination of the goods of value based and policy gradient approaches.

Interaction between 2 agents:

- ▶ an actor: a policy π that acts
- ▶ a critic: a value function that assesses the actions taken by the actor.
→ actor-critic algorithms.

Reinforcement Learning

Actor-critic

- ▶ $\mathcal{L}(\theta) \equiv -Q \log (\pi^\theta)$
→
 $\mathcal{L}(\theta) \equiv -(Q - b) \log (\pi^\theta).$
where b is a function of the (current) state.

Reinforcement Learning

Actor-critic

- ▶ $\mathcal{L}(\theta) \equiv -Q \log (\pi^\theta)$
→
 $\mathcal{L}(\theta) \equiv -(Q - b) \log (\pi^\theta).$
where b is a function of the (current) state.
- ▶ $A(x, a) \equiv Q(x, a) - V(x)$ is the *advantage* function.
→
 $\mathcal{L}_{AC}(\theta) \equiv -(Q - V) \log (\pi^\theta).$

Reinforcement Learning

Actor-critic

- ▶ $\mathcal{L}(\theta) \equiv -Q \log (\pi^\theta)$
→
 $\mathcal{L}(\theta) \equiv -(Q - b) \log (\pi^\theta).$
where b is a function of the (current) state.
- ▶ $A(x, a) \equiv Q(x, a) - V(x)$ is the *advantage* function.
→
 $\mathcal{L}_{AC}(\theta) \equiv -(Q - V) \log (\pi^\theta).$
- ▶ We represent V with an other NN, and we update it with the temporal difference (as in DQN).

Reinforcement Learning

Actor-critic

initialize θ_V and θ the weights for V and π

repeat

 perform K episodes and collect the transitions in T .

 update θ_V by minimizing the MSE of the TD on T

for each episode k **do**

for each time-step t of episode k **do**

$$\hat{Q}_{k,t} \leftarrow \sum_{j=t}^{j=|\tau_k|-1} \gamma^{j-t} r_{k,j}$$

$$\theta \leftarrow \theta + \alpha \gamma^t (\hat{Q}_{k,t} - \hat{V}_{k,t}) \nabla_\theta \log (\pi^\theta(x_{k,t}, a_{k,t}))$$

end for

end for

until ...

Reinforcement Learning

Actor-critic

- ▶ Research on actor-critic algorithms is very very active today.

Reinforcement Learning

Actor-critic

- ▶ Research on actor-critic algorithms is very very active today.
- ▶ Algorithms keep on progressing.

Reinforcement Learning

Actor-critic

- ▶ Research on actor-critic algorithms is very very active today.
- ▶ Algorithms keep on progressing.
- ▶ Many interesting ideas could be detailed:

Reinforcement Learning

Actor-critic

- ▶ Research on actor-critic algorithms is very very active today.
- ▶ Algorithms keep on progressing.
- ▶ Many interesting ideas could be detailed:
 - ▶ how to regularize?

Reinforcement Learning

Actor-critic

- ▶ Research on actor-critic algorithms is very very active today.
- ▶ Algorithms keep on progressing.
- ▶ Many interesting ideas could be detailed:
 - ▶ how to regularize?
 - ▶ how to optimize?

Reinforcement Learning

Actor-critic

- ▶ Research on actor-critic algorithms is very very active today.
- ▶ Algorithms keep on progressing.
- ▶ Many interesting ideas could be detailed:
 - ▶ how to regularize?
 - ▶ how to optimize?
 - ▶ representation issues

Reinforcement Learning

Actor-critic

- ▶ Research on actor-critic algorithms is very very active today.
- ▶ Algorithms keep on progressing.
- ▶ Many interesting ideas could be detailed:
 - ▶ how to regularize?
 - ▶ how to optimize?
 - ▶ representation issues
 - ▶ etc

Reinforcement Learning

Actor-critic

Currently, state-of-the-art algorithms are:

- ▶ DQN when the number of possible actions is small.
- ▶ actor-critics like PPO, SAC.

Reinforcement Learning

Application: TD-Gammon



Early 1990's, a real tour de force.

Reinforcement Learning

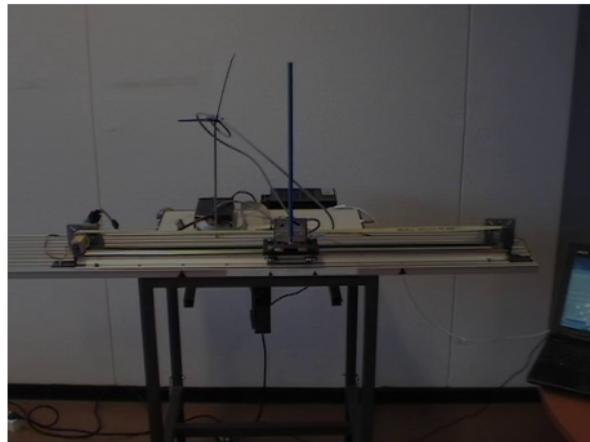
Application: TD-Gammon

- ▶ Backgammon is studied at least since 1974
- ▶ Branching factor: 800
- ▶ TD-Gammon: “successor of NeuroGammon, trained by supervised learning. NeuroGammon won the 1st Computer Olympiad in London in 1989, handily defeating all opponents. Its level of play was that of an intermediate-level human player.” (Source: wikipedia)
- ▶ raw representation of the board position
- ▶ trained with $\text{TD}(\lambda)$ algorithm
- ▶ no knowledge, self-play
- ▶ hand-crafted features
- ▶ 3-plies in v3

Tesauro, Temporal Difference Learning and TD-Gammon, *Communications of the ACM*, 1995

Reinforcement Learning

Application to robotics: a real cartpole

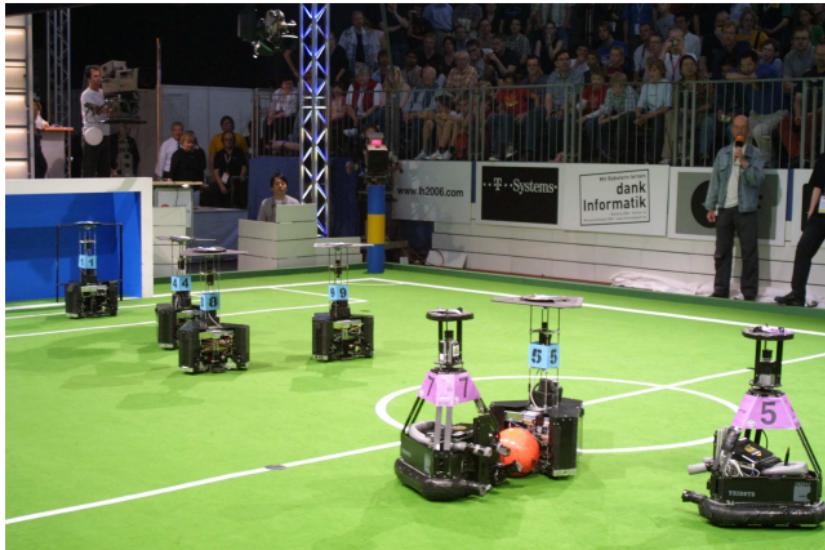


Neural Q-Learning. 2 hidden layers each made of 5 neurons with sigmoid activation function.

Riedmiller, Neural reinforcement learning to swing-up and balance a real pole, *Proc. 2005 IEEE International Conference on Systems, Man and Cybernetics*

Reinforcement Learning

Application to robotics



RL is used only to control the ball and the speed of the robot.

Lauer et al., Cognitive Concepts in Autonomous Soccer Playing Robots, *Cognitive Systems Research*, 11(3), 287:309, September, 2010
(No Deep Learning! only shallow multi-layer perceptron)

Reinforcement Learning

Application: board games

- ▶ Learning to play board games using only the rules of the game.
- ▶ Alpha Go learned to play Go by using games played by humans.
- ▶ Alpha Zero learned to play even better by itself by RL.
- ▶ then other board games (chess, draughts, reversi, ...).
- ▶ then Starcraft II.

Reinforcement Learning

Alpha Zero type of algorithms

- ▶ RL
- ▶ + various tricks to stabilize learning and make it more efficient
- ▶ MCTS as a key component
- ▶ moderately deep network as function approximator

Outro

Some big questions:

- ▶ learning representation
- ▶ generalization in RL
- ▶ time varying environments
- ▶ transfer learning
- ▶ life-long learning
- ▶ risk-aware RL
- ▶ explanation/accountability of the learned behavior
- ▶ hierarchical RL
- ▶ combination of symbolic AI with RL
- ▶ *etc*

Outro

Some big questions:

- ▶ learning representation
- ▶ generalization in RL
- ▶ time varying environments
- ▶ transfer learning
- ▶ life-long learning
- ▶ risk-aware RL
- ▶ explanation/accountability of the learned behavior
- ▶ hierarchical RL
- ▶ combination of symbolic AI with RL
- ▶ *etc*

Advice: there are many many topics to study in RL. RL is not yet mature, a lot is yet to be discovered.

However, to have any chance of success, you need to understand what you are doing, not just run something and hope it will do something: in general it won't.

Reinforcement learning resources

The RL community is open source.

Lots of excellent codes available online.

Some nice RL libraries:

- ▶  **rlberry**: our home-brewed RL library:
<https://github.com/rlberry-py/rlberry>
- ▶ stable-baselines3:
<https://stable-baselines3.readthedocs.io/en/master/>
- ▶ clean-RL: <https://github.com/vwxyzjn/cleanrl>
- ▶ spinning-up RL: <https://spinningup.openai.com/en/latest/>
- ▶ mushroom RL:
<https://mushroomrl.readthedocs.io/en/latest/>
- ▶ ...

Bibliography

- ▶ Bertsekas, A Course in Reinforcement Learning, Athena Scientific, 2023, available on the web for free
<http://web.mit.edu/dimitrib/www/RLbook.html>
- ▶ Lapan, *Practical Deep Reinforcement Learning*, Packt, 2018
- ▶ Puterman, *Markov decision processes*, Wiley, 1994
- ▶ Silver *et al.*, Mastering the game of Go without human knowledge, *Nature*, **550**, 2017
- ▶ Tesauro, Temporal Difference Learning and TD-Gammon, *Communications of the ACM*, 1995

Thanks for your attention!

and congratulations!

and have fun with reinforcement learning!

Credits

All figures are mine, except:

- ▶ slide 3, the backgammon image comes from <https://en.wikipedia.org/wiki/Backgammon>, the go image comes from ?, the chatGPT logo comes from https://upload.wikimedia.org/wikipedia/commons/0/04/ChatGPT_logo.svg, and the data center image comes from <https://www.edgeir.com/energy-monitoring-for-edge-data-centers-20220902>.
- ▶ Slide 30, image comes from <https://www.liveabout.com/craps-dice-control-537459>
- ▶ Slide 63, the figures come from Rémi Munos' website <http://researchers.lille.inria.fr/munos/variable/index.html>.
- ▶ Slide 41, the donkey was designed and made by Sertan Girgin in 2008 while he was a post-doc in my research group, SequeL.
- ▶ Slide 71, the Atari image comes from <https://deepai.org/publication/playing-atari-games-with-deep-reinforcement-learning-and-human-checkpoint-replay>.
- ▶ Slide 82, the Backgammon game image comes from <https://en.wikipedia.org/wiki/Backgammon>.
- ▶ Slide 84, the cartpole image is one frame from a video of M. Riedmiller's Neuroinformatics group at U. Osnabrueck, <https://www.riedmiller.me/subprojects/data-efficient-rl>.
- ▶ Slide 85, the image is one frame from a video of M. Riedmiller's Neuroinformatics group at U. Osnabrueck, <https://www.youtube.com/watch?v=0vNXLVmExyI>.
- ▶  image comes from <https://newquayjunior.net/homework/>.